

Every time an 8(a) subsidiary graduates, Bristol Bay loses the proposals that won its work.

ANCHORAGE, ALASKA 2026-07-28

IN SHORT

Bristol Bay's government services group runs more than 300 million dollars a year in federal contracts across 22 separate 8(a) subsidiaries. Because 8(a) status lasts only nine years, those subsidiaries are always aging out and being replaced. That churn is the model working as designed. What is not by design is that each departure takes proven, winning proposal text with it, because that text lives on the drives of the entity that wrote it and nowhere else.

We looked at how the group wins federal work and where the same effort gets spent twice. The finding is narrow. Across 917 government contractors surveyed this year, a single proposal takes an average of 84 hours, which is two full work weeks. A large share of that goes into past performance, resumes, and compliance answers that a sister company in the same family already wrote and already won on. Nobody can find them, so they get written again.

What we would build is a searchable memory of Bristol Bay's own winning proposals. A writer asks in plain language for what a sister company has already done, and gets back the actual paragraph with the entity and the contract it won attached. They check it and reuse it. Nothing is written for them and nothing is submitted automatically. The rules about who may see which document stay ordinary

software rules, tested and auditable, because a permission boundary is not something to hand to a language model.

We would start with one subsidiary, not all 22. That pilot costs roughly 80,000 to 133,000 dollars over five years, and the conservative case does not pay that back in cash on recovered hours alone. What it does return, inside six to nine months, is a number nobody has today: what a proposal actually costs Bristol Bay. That number decides whether the larger build is worth funding.

One caveat worth reading before the rest. We could not verify your real proposal hours or win rate, because they are not public. Everything modeled here uses an industry average and an assumed 88 dollar hourly rate, both stated plainly so you can correct them. If those assumptions are wrong the numbers move, and we would rather you check them than take them.

01

The library that would make the next bid faster already exists inside the family

It is just scattered across 22 companies that cannot search each other.

Bristol Bay's government services group is 22 subsidiaries employing more than 2,500 people worldwide and executing over 300 million dollars in contracts a year, for the Defense Health Agency, the Air Force, the Army, and the Nuclear Regulatory Commission. Many hold 8(a) certification, which is what lets them win sole-source federal work.

That certification expires after nine years. When it does, the subsidiary graduates and a new one is stood up behind it. This is the model working, and it is a large part of why the group has grown. It also means the family is permanently writing proposals, and permanently losing the ones it wrote.

Each of those 22 entities is a separate legal company with its own drives and its own proposal history. When a writer at one needs a past performance narrative for a Defense Health Agency contract, there is no way to discover that a sister company wrote and won exactly that two years ago. So it gets written from scratch. When that sister company graduates, its version is gone for practical purposes, and usually so is the person who wrote it.

22 companies, one family, no shared memory

The work is already done and already proven. It is only unfindable.

02

The repeated work runs to roughly two work weeks a proposal

We cannot see your real number from outside, so here is the industry's and what we would do about that.

Deltek's 2026 study of 917 government contractors found that developing a single proposal still takes an average of 84 hours. That is two full weeks of one person's time, per bid. The same study found 83 percent of contractors missed opportunities last year simply because they learned about them too late.

That 84 hours is an industry average and not Bristol Bay's number, and we have no way to see yours from the outside. Your group may well be under it. Measuring that is the first thing we would do and it matters more than any projection in this study, which is why the pilot is built to produce it rather than assume it.

What we can say from outside is where repetition is structurally guaranteed rather than merely likely. Past performance narratives, key personnel resumes, and standard compliance answers get written fresh in each entity because there is no shared place

to look. Those are the sections that change least between bids and that a sister company is most likely to have already written well.

03

The goal is not writing faster, it is not writing the same thing twice

What actually changes for a proposal manager on the Tuesday after this goes live.

A solicitation lands on a Tuesday. Here is that same Tuesday before and after, step by step.

The same proposal, before and after, step by step

TODAY	AFTER THE PILOT
1 Solicitation lands. The writer tries to remember which of the 22 entities has bid this agency before.	Solicitation lands. The writer asks the question in plain language.
2 Emails around the family asking whether anyone has something similar, then waits on replies.	Nothing to wait on. The question is answered in about a minute.
3 Digs through shared drives and old submissions for the file, if a helpful reply ever comes.	Passages come back with the entity and the contract they came from attached.
4 Rewrites past performance, resumes and standing compliance answers close to from scratch.	Starts from proven winning text and edits down to what still applies.
5 Chases down which contract each reused claim actually came from before it can be cited.	The citation is already attached to the passage.
6 A cleared person reviews the draft for controlled and export-restricted content.	A cleared person reviews the draft for controlled and export-restricted content.

The last row is deliberately identical. Compliance review is the one step we would not touch, because a human owning that call is the control, and moving it into software is the change that eventually costs you a contract.

The sections that get reused are the ones nobody enjoys writing and reviewers scrutinise least: past performance, resumes, standing compliance answers. The time that frees goes to technical approach and price, which are the parts that actually win. This does not make anyone a faster writer. It stops the family paying twice for the same paragraph, and it stops the knowledge walking out the door when a subsidiary graduates.


A searchable memory of your own winning proposals

Four moving parts, three of which are ordinary software.

The four pieces, in plain terms, before the diagram. A collector that copies proposal documents out of each subsidiary's existing storage, read only. A librarian step where a cleared person marks which documents are controlled and who may see them. A search layer that finds the right passage when someone asks a question in ordinary English. And a check that refuses to show any answer it cannot trace back to a real document.

The writer's side is a single search box. They type what they need the way they would ask a colleague, and they get back real passages from real winning bids with the source attached. They can open the original. Nothing is generated for them, nothing is invented, and nothing is submitted anywhere.

The part that decides who may see which document is deliberately not clever. It is a rule, written down, tested exhaustively, and sitting outside the search entirely. Controlled unclassified information carries real obligations under CMMC and DFARS, and technical data under export control cannot be reachable by non US persons. A rule can be audited line by line and proved correct. Judgement cannot. So the boundary stays a rule, and it filters documents before the search ever sees them.



Read it left to right. Nothing reaches the search until the permission rules have already removed every document this particular person is not cleared to open, so a mistake in the search cannot become a disclosure.

A working demonstration, built on your own facts. It is a demonstration and it is honest about that, and it does nothing this study did not scope. [Open it full screen.](#)

There is exactly one place in this build where the technology genuinely earns its keep, and that is understanding a question asked in ordinary English and matching it to a passage written in completely different words. Keyword search does not do that. Everything else here is deliberately conventional and we would rather say so than dress

it up. Checking whether a subsidiary is eligible for a set-aside is a lookup with a right answer. Deciding who may open a controlled document is a permissions rule. Both stay exactly that.

Most of this is bought rather than built. The secure environment, the search infrastructure, and the identity system are established products and we would use them rather than reinvent them. What is worth building is the part nobody sells you: the rules encoding which of 22 entities may see what, the proof that every answer traces to a real document, and the unglamorous work of getting scattered archives into one place.

05

What it costs, and what it honestly returns

One subsidiary, five years, and a conservative case that does not clear on cash alone.

This prices a pilot on one subsidiary's archive, not a rollout across all 22. The largest cost is not software. It is the unglamorous work of collecting scattered documents and having a cleared person mark what is controlled, and that is also the part we cannot size accurately from outside.

Every figure below assumes a fully loaded proposal writer at 88 dollars an hour and the industry's 84 hours per proposal. Both are stated so you can replace them with your own.

Modeled over five years for a single subsidiary pilot

	Conservative	Most likely	Aggressive
📌 Proposals in the pilot slice, per year	20	30	45
📌 Share of prep hours recovered	8%	12%	18%
📌 Capacity returned per year	\$11,800	\$26,600	\$59,900
📌 Five year cost, all in	\$133,200	\$100,300	\$80,500
📌 Share of that cost recovered	39%	119%	342%
📌 Cash payback	not within five years	about month 45	about month 14

Recovered hours are capacity, not cash. They become money only if the freed time defers a hire, and we have not assumed it does. Earlier discovery of missed bids is deliberately left out of these figures, because win rate is the number most easily inflated and we did not want it carrying the case.

📌 modelled from stated drivers 📌 assumed, and named as such

The conservative case does not pay back in cash

On recovered hours alone it returns about 39 percent of its cost over five years, and we are not going to present that as a win. What the pilot does return, inside six to nine months, is your own measured cost per proposal and your own reuse rate. Nobody has those numbers today, and they are what decides whether the build across all 22 subsidiaries is worth funding or worth dropping.

The honest backdrop is that most of this fails. MIT's work puts about 95 percent of enterprise AI pilots at no measurable profit impact, and RAND finds more than 80 percent of AI projects fail outright. Government contracting looks the same. Adoption

jumped from 45 to about 90 percent this year while average margins fell from 20 to 17 percent, and only around 5 percent of contractors describe their capability as fully developed. Adoption is plainly not the differentiator. What separates the few that work is that they are narrow, they run on the company's own material, a person stays accountable for the output, and somebody measures the result. That is the shape of this one, and it is why we scoped it to one subsidiary and one job rather than something more impressive.

For any of this to mean anything, one person has to own the number. We would want a named proposal manager in the government services group who signs off the baseline at the start and reviews actual against baseline each quarter. That review is the gate. If the numbers do not move, the honest answer is to stop, and we would say so.

06

How we would start, and what would stop us

Nothing funds the next phase except measured results from the last one.

The first two to four weeks are not a build at all. They are a survey of one subsidiary's archive, answering the question that actually determines cost: how much of this material is machine readable, how much is duplicated, and how long it takes a cleared person to mark a single document. That produces a per document number, and that number is what a rollout across 22 entities should be priced against. If it comes back ugly, you have spent very little to find that out.

Phase	What runs	How we would know it worked
Now	Survey one subsidiary's proposal archive and measure the real cost to onboard a document	A per document cost, and a clear answer on whether a wider rollout is weeks or quarters of work
	Build the search over that one archive, with the permission rules in front of it	Writers find a usable prior passage for most real questions, and controlled documents never surface to anyone not cleared for them
	Measure what a proposal actually costs you today	A real baseline for hours per proposal and reuse rate, replacing the industry averages used in this study
Next	Extend to the rest of the government services family	The same accuracy and permission results hold on a larger and messier archive
	Draft help for the compliance matrix, with a person signing every section	Turnaround improves against the baseline, with no unreviewed text reaching a proposal
Later	The same library for the construction group's bids	Reuse rate on construction proposals, measured the same way

Each phase is released by the previous phase's measured numbers rather than by a projection. If the survey says the archive is not ready, the right decision is to fix that first, or to stop.

What we would need from you is modest, but it is not nothing. Access to one subsidiary's proposal archive in an environment your security people approve. A proposal manager for a few hours to tell us what a good answer looks like. And a decision from whoever owns compliance about which documents are controlled, because that classification does not exist yet and it determines what the library may hold.

What would make us wrong

The three most likely ways this disappoints you, and what we would do about each.

The archive turns out to be the whole project. If those 22 archives are mostly scanned images, heavily duplicated, and unclassified for control purposes, then collecting and marking them is the real cost and the search is a footnote. This is the single most likely way our numbers are wrong. It is also why the first step is a survey rather than a build, so you learn it in a few weeks rather than a year.

People may not use it. This is the most common reason pilots like this show nothing, and better technology does not fix it. A writer who does not trust the tool goes back to emailing colleagues, and the recovered hours never appear. We would instrument use from the first day and treat low use as a result worth reporting rather than a problem to hide.

Search quality degrades as an archive grows. More documents makes retrieval harder rather than easier, and keeping it accurate is ongoing work rather than a one time setup. We would rather say that now than discover it with you in month eight.

And the one already said plainly above. On recovered hours alone the conservative case does not pay back inside five years. It is honest as a pilot that buys you a measured number. If what you need is a build that pays for itself on hours saved, this is not that one, and you should not fund it on that basis.

08

The next step is small

This study is the free version of how we work. There is nothing attached to it and no follow up sequence behind it.

If it is useful, the next step is a reply rather than a commitment. The most valuable thing we could do together is the archive survey, because it turns the largest unknown in this

document into a number you own. And if we have read something wrong about how your group actually operates, that is worth ten minutes too. We would genuinely want to know.

What we checked

Every number in this study traces to a page we read. If we have any of it wrong, that is worth knowing too.

- 1** Government services group scale, 22 subsidiaries, over 2,500 personnel, more than 300 million dollars in contracts a year, and the agencies served
<https://eiseservices-llc.com/about/>
- 2** Deltek Clarity 2026, 917 contractors surveyed, 84 hours per proposal, 83 percent missed opportunities, adoption 45 to 90 percent, margins 20 to 17 percent, about 5 percent fully developed
<https://www.govconwire.com/articles/clarity-report-2026-deltek-govcon-kevin-plexico>
- 3** Bristol Bay Native Corporation, ANCSA regional corporation incorporated June 13 1972, about 9,900 shareholders as of 2015
https://en.wikipedia.org/wiki/Bristol_Bay_Native_Corporation
- 4** Strongest operating year in company history, 3.2 billion dollars revenue, 202 million earnings, 84 million net income
<https://www.akbizmag.com/industry/alaska-native/regionalroundup2025/>
- 5** Bristol Bay Construction Holdings and its five construction brands
<https://www.bbnc.net/business/bristol-bay-construction-holdings/>
- 6** Koniag Government Services, the CASPER platform and Koniag AI Solutions, with vendor published figures including a 75 percent reduction in contract requirements analysis time and 29.3 million dollars in savings
<https://www.koniag-gs.com/ai-investments-deliver-powerful-results-for-koniag-government-services-and-federal-customers/>

- 7** Alutiiq bid and capture automation on Salesforce Government Cloud, reported at 100 percent adoption after one year
<https://www.salesforce.com/customer-success-stories/alutiiq/>
- 8** ASRC Federal named AI and machine learning deployments
<https://www.asrcfederal.com/ai/>
- 9** Compliance limits on AI touching controlled unclassified information, covering CMMC, NIST 800-171, FedRAMP, export control, and False Claims Act exposure
<https://www.kiteworks.com/regulatory-compliance/ai-compliance-federal-contractors/>
- 10** SPATHE Systems, a comparable 8(a) contractor, vendor published proposal turnaround results
<https://www.govdash.com/blog/from-3-weeks-to-24-hours-how-spathe-systems-4x-d-rfp-output-with-govdash>
- 11** MIT Project NANDA, about 95 percent of enterprise generative AI pilots show no measurable profit impact
<https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/>
- 12** RAND report RRA2680-1, more than 80 percent of AI projects fail
https://www.rand.org/pubs/research_reports/RRA2680-1.html
- 13** Bristol Bay Services, the government services group and its leadership
<https://www.bbnc.net/business/bristol-bay-services/>