

Bristol Bay's most valuable proposal asset is the proposals it has already won, and today that memory leaks out the door every time an 8(a) subsidiary graduates.

The government-services engine runs more than 300 million dollars a year across 22 separate 8(a) subsidiaries, and every one of them writes past performance, resumes, and compliance answers from scratch in its own silo. Because 8(a) certification lasts only nine years, older firms graduate and new ones are stood up constantly, so the group is on a perpetual proposal treadmill and each graduation walks proven, winning text out the door. The library that would make each new bid faster and stronger already exists inside the family. It is just scattered across 22 legally separate entities, trapped in PDFs and per-sub drives, never unified. The opportunity is not to buy an AI writer. It is to turn your own winning work into a governed, searchable memory before more of it is lost.

Anc

Anchorage, Alaska

2026-07-23

Researched and drafted by Alaska AI's agent team, reviewed by a human before it reached you

WHAT WE SAW

The homework

THE BUSINESS

Bristol Bay Native Corporation is one of the thirteen Alaska Native regional corporations created under ANCSA, incorporated in Alaska on June 13, 1972, and headquartered in Anchorage. It is a diversified holding company owning subsidiaries across federal government services, construction, industrial and petroleum services, seafood, tourism, and a newer financial-services line, and it returns profit to roughly 9,900 Alaska Native shareholders as of 2015, a base now opening to descendants. Last fiscal year it reported its strongest operating performance in history, 3.2 billion dollars in revenue, 202 million in earnings, and 84 million in net income. Its largest engine is federal contracting through a family of 22 government-services subsidiaries, many of them SBA-certified 8(a) small disadvantaged businesses, together employing over 2,500 personnel worldwide and executing more than 300 million dollars in contracts per year for agencies including the Defense Health Agency, the US Air Force, the US Army, and the Nuclear Regulatory Commission.

THE REAL PAIN

Proposal, capture, and 8(a) compliance work is spread across 22 separate government-services subsidiaries, each a distinct contracting entity. As older 8(a) firms graduate the nine-year program, new ones are stood up and have to get onto contracts, so solicitation response and compliance are a permanent high-volume load, not a one-time build. The industry backdrop makes the cost concrete. Across 917 government contractors surveyed for the 2026 Deltek Clarity study, proposal development still averages 84 hours per proposal, and 83 percent of contractors missed opportunities in 2025 because they discovered them too late. Every entity that graduates takes its winning past performance with it.

Source: <https://eisservices-llc.com/about/>

THE FIELD AROUND THEM

- **Koniag Government Services** A peer Alaska Native corporation and the

AI IN THIS INDUSTRY NOW

- AI adoption among government contractors jumped from 45 percent in

most aggressive AI mover in the group. It built a proprietary enterprise language-model platform called CASPER inside a federal-compliant environment and launched a dedicated Koniag AI Solutions subsidiary. Its own published figures (vendor-reported, not independently audited) claim a 75 percent reduction in time for contract requirements analysis, a 50 percent reduction in human time on language determination across more than 10,000 documents, and 29.3 million dollars in savings from a procurement tool that cut a process from one hour to five minutes. [source](#)

- **Alutiiq (Afognak)** A smaller peer that modernized bid and capture with automation, not AI: a bid-and-capture application on FedRAMP-authorized Salesforce Government Cloud with a full pipeline view and automated status reporting, reaching 100 percent user adoption after one year. Notable because it shows the capture-workflow bar is already being cleared by peers with plain automation. [source](#)
- **ASRC Federal (Arctic Slope)** Runs named AI and machine-learning deployments, an analytics platform for a NOAA financial system and an anomaly-detection tool for GOES-R satellite data. This is mission-delivery AI sold to federal customers, not internal capture or proposal work, so it

2025 to about 90 percent in 2026, and it is landing specifically in business development, opportunity identification, and proposal automation (Deltek Clarity 2026, 917 contractors, an authoritative industry survey). [source](#)

- A comparable 8(a) contractor, SPATHE Systems, reports (vendor-published, not independently audited) that an AI capture tool drafting from its own past performance cut proposal turnaround from 2 to 3 weeks down to 24 hours and took its no-notice RFP output from 2 per month to 7 to 8 per month with no added proposal staff. The pattern that works is training on the firm's own library with a human in the loop, not a general chatbot. [source](#)
- The hard limit is real and it shapes the whole design. AI that touches controlled unclassified information has to meet CMMC 2.0, NIST 800-171, and DFARS controls, many commercial AI tools are not FedRAMP-authorized, ITAR-controlled technical data reachable by non-US persons can be an unauthorized export, and certifying CMMC while running ungoverned AI over that data can create False Claims Act exposure. That is exactly why this build keeps human review of every compliance-graded document in the loop. [source](#)

points at a different use than the one here. [source](#)

THE OPPORTUNITY

When a new solicitation lands anywhere in the 22-subsidary family, help that subsidiary respond faster and better by drawing on the whole group's own past-performance and proposal memory, with a compliance reviewer always in the loop, so the group wins more of the federal work it is eligible for without adding a proposal team for every new entity.

It scores highest on the honest test of importance high, satisfaction low. Importance is structural, not seasonal, because the 8(a) graduate-and-replace churn never stops. Satisfaction is near the floor, because your own government-services pages show no AI or automation while peers publish results. And it fits where the corporation is going, protecting and growing the 300-million-dollar engine at its strongest year ever.

WHAT THEY USE TODAY

Today the workaround is a manual scavenger hunt: a writer emails around 22 separate entities, digs through per-sub drives and old PDFs, and often rewrites from scratch what a sister company already wrote and won on.


WHY A CHANGE IS REALISTIC

The push is strong (a perpetual proposal load and a visible peer lead), and the pull is strong (proven reuse and a first-mover position no peer has claimed). The force that actually decides it is anxiety, the CUI, CMMC, and ITAR rules and the base rate that most AI pilots show no measurable return. That is why this is built governance first, so the anxiety becomes the design spec rather than a reason to wait.

THE RECOMMENDED BUILD

Capture Memory

A governed retrieval layer over Bristol Bay's own past-performance narratives, prior proposals, resumes, and compliance-matrix answers across the government-services subsidiaries. A proposal writer asks in plain language, for example who has Defense Health Agency medical-staffing past performance we can reuse, and gets back the actual prior text with the winning entity and contract cited beside it. The writer verifies the source and reuses it. Nothing is generated from thin air, nothing is auto-submitted, and every answer traces to a real document. A companion piece, Opportunity Radar, watches public solicitation feeds and flags fitting opportunities and deadlines early.



■ Person ■ External ■ Their system ■ We build ■ Data

Deterministic access control sits outside the model and filters before retrieval, so the model only ever sees and cites text the writer is cleared for. Nothing is auto-submitted. Opportunity Radar is a parallel deterministic lookup, not an AI product.

WHERE AI FITS, AND WHERE IT DOES NOT

AI genuinely fits exactly one part of this, understanding a plain-language question and matching it to relevant prior text across entities where keyword search fails. That is language understanding, which is a real strength. Everything

BUILD VS BUY

Buy the commodities, the FedRAMP-authorized environment, the compliant model endpoint, the embeddings and vector store, and identity. Build only your edge, the deterministic access-control layer, the citation and traceability that makes every answer checkable, and the

around it is deliberately not AI. The rule that decides who may see which document under CUI, ITAR, and 8(a) boundaries is deterministic access control that sits outside the model and is tested exhaustively, because a permission boundary is never something to hand to a language model. The NAICS and set-aside eligibility matching in Opportunity Radar is a lookup with a right answer, so it stays a rule, not a model. Saying where we would not use AI is the point, it is what earns the part where AI genuinely wins.

cross-subsidiary corpus that unifies memory no vendor can assemble for you.

THE PLAN

What we would build, scoped

The government-services group writes proposals against the same kinds of past performance, resumes, and compliance answers, from scratch, in 22 separate silos, and each 8(a) graduation walks proven winning text out the door. The reusable library already exists inside the family, it is just not searchable across entities.

GOALS

- Turn the group's scattered past performance, prior proposals, resumes, and compliance answers into one governed, searchable memory where a writer asks in plain language and gets sourced, reusable prior text.
- Make every answer traceable to a real retrieved source, so a writer verifies before reuse and nothing untraceable can reach a federal proposal.
- Surface qualifying opportunities earlier with a deterministic eligibility and deadline lookup.
- Measure Bristol Bay's real proposal-hour, reuse, and win-rate baselines during the pilot so the decision to fund the next phase rests on your own data, not a vendor promise.

NON-GOALS, ON PURPOSE

- No autonomous proposal drafter in phase one. Capture Memory is generation-light retrieval only. (An end-to-end autonomous chain runs eight to ten steps at roughly 90 percent reliability each, landing near 35 to 60 percent overall, which is unacceptable for compliance-graded documents where a wrong claim is real legal exposure. Drafting is sequenced to a later, human-checkpointed phase once a governed library exists.)
- No letting a model decide who may see which document. (CUI and ITAR access is deterministic rules that can be tested exhaustively and must sit outside the model, so a boundary failure can never be a model failure.)
- No AI on the NAICS and set-aside eligibility matching. (These are lookups with a right answer, so a model only adds hallucination risk and no accuracy.)
- Nothing is auto-submitted, and the model never files or sends a proposal.

(A writer owns and verifies every reuse, and human ownership of compliance-graded text is mandatory.)

- No touching the number-generating finance seat in phase one. (Bristol Bay is standing up its own centralized financial planning, and inaccurate AI forecasts are a top peer complaint, so the pilot stays on retrieval of existing text, not generation of numbers.)

HOW WE WOULD MEASURE SUCCESS

METRIC	BASELINE	TARGET	BY WHEN
Retrieval precision at k on a proposal-manager-built evaluation set (top results are relevant, reusable prior text)	0, no retrieval system exists today (assumed, measured at first evaluation run)	at least 80 percent of evaluation queries return a correct, reusable source in the top three results	by end of the pilot build, before writer rollout
Citation faithfulness and cross-boundary leakage (every returned sentence traces to a real source, zero controlled text surfaced to a non-cleared identity on a red-team permission set)	must be perfect, this is a governance gate not a dial	100 percent of returned text traceable, and 0 leakage events across the full permission-boundary test set	continuous, gating go-live and re-run every release
Proposal-hours per response on pilot-slice bids	84 hours per proposal (assumed from the 2026 industry average, measured on the pilot slice in the first weeks before tool exposure)	a credible reduction on reused-section prep, conservatively modeled at about 8 percent of prep hours recovered via found-and-cited prior text	measured across pilot-slice proposals over the 90-day pilot
Reuse rate (share of eligible proposal sections drafted from a cited prior-	near 0 today, prior text is not cross-searchable	at least 30 percent of eligible sections sourced from a cited reuse	by end of the 90-day pilot

METRIC	BASELINE	TARGET	BY WHEN
win source rather than from scratch)	(measured from writer logs at pilot start)		
PHASE ONE, MUST HAVE • Governed retrieval over a bounded corpus slice, the shared government-services library or one high-volume subsidiary, plain-language query in, sourced reusable prior text out with the winning entity and contract cited. • Deterministic access control (CUI, ITAR, 8(a) permissioning) sitting entirely outside the model, tested against a red-team permission set that must return not cleared. • Citation-first behavior, every answer traces to a source, low-confidence retrievals route to a human, and an honest I do not have that over any fabrication. • A separate screening instance that flags untraceable text before it reaches a writer, so the retrieval model never polices itself. • Opportunity Radar, a deterministic eligibility and deadline lookup over public feeds, with an optional fit-score clearly labeled as a suggestion. • Baseline instrumentation from day one, proposal-hours, reuse rate,		DELIBERATELY LATER • A human-checkpointed compliance-matrix and first-draft assist (Shred-to-Draft), funded only after the governed library is proven. • Extending the library across all 22 subsidiaries plus the construction group once the bounded slice proves value and governance. • A construction-group bid-and-capture extension. • Back-office uses (legal and contract administration, payroll compliance, shareholder-season deflection), each earning its own gate later.	

retrieval precision, and opportunities surfaced early.

WHAT WE WOULD NEED FROM YOU

- Access to a bounded corpus slice, the shared government-services library or one high-volume subsidiary's prior proposals, resumes, and compliance answers, in an agreed governed environment.
- A proposal manager's time to co-build the evaluation set, a few dozen real proposal question-and-answer pairs plus a permission-boundary test set.
- Per-document CUI and ITAR classification decisions from your compliance owner, this classification is not done today and it gates what the retrieval layer may index.
- Sign-off on the governed environment and its CMMC posture, and confirmation of who counts as cleared.
- A named pilot sponsor in the government-services group and a librarian or proposal lead to own unresolved queries.
- Your true proposal-hours and win-rate baselines on the pilot slice, or the go-ahead to measure them before tool exposure.

RISKS, AND HOW WE WOULD HANDLE THEM

RISK	MITIGATION
The corpus is more scattered, less deduplicated, and less freshly permissioned than expected, and assembly and classification eat the timeline. This gap is the project.	Phase one is bounded to one slice precisely so assembly is finite, per-document classification is a stated dependency owned by your compliance team, and we size the pilot around the real state of the slice, not an assumed clean corpus. A two to four week corpus-readiness spike on one subsidiary returns a per-document cost-to-onboard before any group-wide dollar is committed.
A hallucinated or mis-cited past-performance claim reaches a federal proposal.	Retrieval first with a citation on every sentence, a separate screening instance flagging untraceable text, mandatory human verification before reuse, and nothing auto-submitted.

RISK	MITIGATION
Controlled text is surfaced to a non-cleared person.	Access control is deterministic, sits outside the model, and is validated against a red-team permission set with a zero-leakage gate before go-live.
The pilot underperforms the hour-reduction target, reuse turns out thin, or writers do not adopt.	The pilot is framed so its return is both modest reuse capacity and decision data. Even if hours are not cut, the measured baselines are a real, fundable output that justifies or stops the next spend, and we will say plainly if the honest read is do not expand.

OPEN QUESTIONS WE WOULD RESOLVE TOGETHER

- Which bounded slice does phase one index, the shared government-services library or one named high-volume subsidiary? (Bristol Bay Services leadership with the pilot sponsor)
- What is Bristol Bay's true proposal-hours-per-response and win-rate baseline on the pilot slice today, versus the 84-hour industry assumption? (Pilot-slice proposal manager, measured jointly in the first weeks)
- Where does per-document CUI and ITAR classification stand, and who signs each decision? (Bristol Bay compliance and security lead)
- What governed environment and CMMC posture must the pilot run in, and which identities count as cleared? (Bristol Bay IT security and compliance owner)

THE NUMBERS

Honest ROI, not a hero number

WHAT IT COSTS TO RUN, FIVE YEAR VIEW

This is a five-year picture for a bounded phase-one pilot, not a group-wide rollout. The biggest and most underestimated line is implementation, and inside it the hidden project is corpus assembly, deduplication, per-document CUI and ITAR classification, and permissioning. Bounding the pilot to one slice is what keeps that line in the low-to-mid five figures instead of six. All modeled labor value uses an assumed fully-loaded proposal-writer rate of 88 dollars per hour, stated here so every dollar figure below is auditable. Run cost is retrieval-heavy and generation-light, so inference is low, and it carries monitoring, a standing evaluation set, mandatory human review, and periodic model refresh. A 15 to 20 percent contingency reserve sits on top, because most AI-pilot failure is organizational, scope creep, thin training, and no owner, not the software. Training is carried as real dollars, never assumed included. Modeled five-year total cost of ownership lands near 133,000 dollars in the conservative case, near 100,000 in the most-likely case, and near 80,000 in the aggressive case, the differences driven by the size of the corpus-assembly effort.

WHERE THE VALUE COMES FROM

- **Governed reuse of prior proposal text, found and cited instead of rewritten from scratch** **CAPACITY** Anchored to the 84-hour-per-proposal industry figure used as an assumed baseline until Bristol Bay's own number is measured in the pilot, and valued at an assumed fully-loaded proposal-writer rate of 88 dollars per hour (a stated assumption, confirmable in a five-minute check with your finance team). This is capacity, not cash. It converts to cash only through a defer-hire test, deferring or avoiding a proposal-coordinator hire as 8(a) bid volume rises. If that does not happen, we book it as zero.
- **Earlier solicitation discovery via the deterministic Opportunity Radar** **CAPACITY** Against the 83 percent of contractors who missed opportunities in 2025 from late discovery, recovering even one winnable bid a year is a large upside. It is deliberately kept out of the cash math and named as capacity, because win-rate lift is exactly the number that inflates a case.

- **Measured baselines that gate the next-phase spend** **CAPACITY** The pilot instruments Bristol Bay's real proposal-hour, reuse, and win-rate baselines. That measured data is the actual conservative return, it makes the far larger next-phase commitment fundable or killable on your own evidence rather than a vendor claim.

CONSERVATIVE

20 modeled pursuits in the bounded slice, an 8 percent reuse trim on 84 hours each, a slow 40 percent first-year ramp, and full padded cost. Annual recovered capacity is about 11,800 dollars, and over five years that is about 52,000 dollars against a total cost near 133,000, so on reuse hours alone it recovers about 39 percent of cost and does not reach cash payback inside five years. We say that plainly. It clears the bar as decision data, the pilot cost is a small fraction of the next-phase spend it de-risks, and it produces your measured baselines inside the roughly six to nine month pilot window.

MOST LIKELY

30 pursuits, a 12 percent reuse trim, a 50 percent first-year ramp, and leaner but still padded cost. Annual recovered capacity is about 26,600 dollars, about 120,000 over five years against a total cost near 100,000, recovering about 119 percent of cost, with cash payback around month 45 if the freed capacity is converted through the defer-hire test. A defensible outcome, not a slam dunk.

AGGRESSIVE

45 pursuits, an 18 percent reuse trim, a 60 percent first-year ramp, and the lowest defensible cost. Annual recovered capacity is about 59,900 dollars, about 275,000 over five years against a total cost near 80,000, recovering about 342 percent of cost, with cash payback around month 14. This is the upside case, labeled as such, and nothing in the recommendation leans on it.

The conservative case still clears the bar. That is the case we would stand behind.

Payback as decision data, the pilot pays for itself the moment it produces Bristol Bay's own measured baselines, inside the roughly six to nine month pilot window. Converted to cash through a deferred proposal-coordinator hire, payback runs from about 14 months in the

aggressive case to about 45 months in the most-likely case. On reuse hours alone the conservative case does not reach cash payback inside five years, and we say so.

The honest base rate

The outside view is unforgiving and we state it first. Most generative-AI pilots show no measurable profit impact, the MIT figure is about 95 percent and RAND finds more than 80 percent of AI projects fail. Government contracting mirrors it, about 90 percent adoption in 2026 yet average margins fell from 20 to 17 percent and only about 5 percent of contractors call their AI fully developed. Adoption is not the differentiator, integration and governance are. This build is designed for the small share that pays because it is workflow-embedded, trained on your own past-performance library, human-owned and generation-light with every answer traced to a source, and measured with a named owner tracking actual against baseline. The peer numbers in this study are vendor-published claims, not your guaranteed result.

Who owns the number The government-services group (Bristol Bay Services) owns the number, with a named proposal manager accountable day to day for co-building the evaluation set and signing off the measured baselines. At go-live and quarterly, actual proposal hours and reuse rate on the pilot slice are compared against the pilot baseline, and that variance report is the go or no-go gate for the next-phase spend.

THE ROADMAP

By confidence, not by calendar

NOW

Capture Memory walking skeleton over one bounded corpus slice, running end-to-end, with deterministic access control outside the model and nothing auto-submitted.

Retrieval precision at k and citation faithfulness on a proposal-manager-built set, and zero cross-boundary leakage on a red-team permission set. Ship bar is faithfulness at threshold with zero permission leaks.

Opportunity Radar companion, a deterministic eligibility and deadline lookup over public feeds with an optional labeled fit-score.

In-window, in-eligibility solicitations surfaced per week that the group was not already tracking. Falsifiable bar, recover at least one winnable, previously missed solicitation over the pilot.

Decision-data instrumentation baked into the pilot, baseline proposal-hours, win rate, and reuse capacity recovered.

A measured proposal-hour and win-rate baseline captured for the pilot slice, the gating return even if capacity is never converted to cash.

NEXT

Shred-to-Draft, a human-checkpointed compliance-matrix and first-draft assist grounded in the now-live Capture Memory, never an autonomous agent.

Proposal draft turnaround and response throughput against the measured baseline, with zero un-owned compliance-graded text reaching a proposal.

Extend Capture Memory to the construction group's bid-and-capture, the same governed layer on its own past performance.

Retrieval faithfulness and permission-boundary integrity held on the new corpus at the same bar as the pilot.

LATER

Centralized legal and contract-administration assist, governed retrieval with a human owning every compliance-graded output.

Contract and clause retrieval faithfulness with human ownership, and adoption by the centralized legal function.

Payroll operations and compliance automation across the multi-entity workforce.

Form-review and data-entry hours reduced and compliance error rate, against a measured baseline.

Shareholder-season deflection and FAQ assist, tested against a plain-heuristic baseline before any voice build.

Share of board-season inquiry volume deflected versus the heuristic, build proceeds only if it beats the heuristic.

Each phase is earned by the last phase's measured results. Nothing funds on projection. The pilot is the smallest funded slice, the next phase is released only when Capture Memory is live with citation faithfulness at threshold,

zero leakage, and a captured baseline, and the later lane is scoped only after the draft assist proves the pattern generalizes.

The honest part, what could make this not work

- The corpus is the project. If your past-performance library turns out more scattered, less machine-readable, or less classified than a one-subsidiary spike suggests, assembly and classification, not the AI, become the real cost and timeline. We measure that on one subsidiary before any group-wide dollar is committed.
 - Retrieval degrades on noisy corpora and as they grow. The fix is better retrieval, hybrid search, and reranking, not more documents, and it is real ongoing work.
 - Adoption is the most common reason AI pilots show no impact. If writers do not trust or use the tool, the capacity never materializes, so we instrument usage from day one and treat weak adoption as a measured signal, not a hidden failure.
 - The governance boundary is unforgiving. A single controlled document surfaced to the wrong person is a serious event, which is why access control is deterministic, tested exhaustively, and never handed to the model.
 - On reuse hours alone, the conservative case does not pay back in cash inside five years. It is honest as a decision-data pilot, and if the measured baselines say do not expand, that is a real and useful answer too.
-

WHAT THIS IS

This is the free version of how we work

This study is the free version of how we work, a real Field Study, not a sales deck. If it is useful, the small next step is a reply, not a commitment, and we can walk the corpus-readiness spike that turns the biggest unknown into a number. The full picture of how we work is at alaskaaihq.com/services.

The full picture lives at alaskaaihq.com/services.

SOURCES

1. BBNC incorporated June 13 1972 as an ANCSA regional corporation, about 9,900 shareholders https://en.wikipedia.org/wiki/Bristol_Bay_Native_Corporation
2. Strongest year in history, 3.2B revenue, 202M earnings, 84M net income, descendant enrollment, financial-services expansion <https://www.akbizmag.com/industry/alaska-native/regionalroundup2025/>
3. 22 government-services subsidiaries, over 2,500 personnel worldwide, more than 300M in contracts per year, agencies served <https://eiservices-llc.com/about/>
4. Bristol Bay Construction Holdings, five construction brands <https://www.bbnc.net/business/bristol-bay-construction-holdings/>
5. February 2026 senior corporate-infrastructure hires in HR, finance, and legal <https://www.akbizmag.com/industry/alaska-native/bristol-bay-native-corporation-new-hires/>
6. Koniag CASPER and Koniag AI Solutions, vendor-published metrics <https://www.koniag-gs.com/ai-investments-deliver-powerful-results-for-koniag-government-services-and-federal-customers/>
7. Alutiiq bid-and-capture automation on Salesforce Government Cloud, 100 percent adoption <https://www.salesforce.com/customer-success-stories/alutiiq/>
8. ASRC Federal named AI and ML deployments <https://www.asrcfederal.com/ai/>
9. Deltek Clarity 2026, 84 hours per proposal, 83 percent missed opportunities, adoption 45 to 90 percent, margins 20 to 17 percent, 5 percent fully developed <https://www.govconwire.com/articles/clarity-report-2026-deltek-govcon-kevin-plexico>
10. GovDash and SPATHE Systems, vendor-published proposal outcomes <https://www.govdash.com/blog/from-3-weeks-to-24-hours-how-spathe-systems-4x-d-rfp-output-with-govdash>
11. GovCon AI compliance limits, CMMC, NIST 800-171, FedRAMP, ITAR, False Claims Act exposure <https://www.kiteworks.com/regulatory-compliance/ai-compliance-federal-contractors/>
12. Bristol Bay Services government-services group leadership <https://www.bbnc.net/business/bristol-bay-services/>
13. MIT Project NANDA, about 95 percent of enterprise generative-AI pilots show no measurable profit impact <https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/>
14. RAND (RRA2680-1), more than 80 percent of AI projects fail https://www.rand.org/pubs/research_reports/RRA2680-1.html

Alaska AI. This study was prepared for Bristol Bay Native Corporation and is not for redistribution.