

Aurora season opens in sixteen days and your two pages disagree about the date.

FAIRBANKS, ALASKA 5 AUGUST 2026

IN SHORT

You told us the shape of this yourselves. Your jobs page says you more than double hotel staffing for summer to make sure guests have caring, personal experiences at your properties, and the roster you hire starts with Seasonal Front Desk Clerks. Those are the people standing between a guest and every answer, at three separately addressed properties, on desks you staff around the clock.

The answers are already good. Check-in at 4PM and check-out at 11AM, 24 hours notice on cancellations, no late check-outs from mid-May through mid-September, a courtesy van running 24 hours that excludes tour groups. What that content never does is tell a clerk which answer belongs to the building they are standing in, on the date they are standing there.

Properties your size and shape are getting real measured results from this, and further down we name them, say who measured each one, and say plainly where the evidence is weaker than it looks.

The first thing we would build for you has no AI in it. A set of cards, one question and one agreed answer each, tagged to a property and a date range, with a source, an owner and a review date. A season that flips because the calendar says so rather

than because someone remembered. Plus a sheet of paper we think will be the most-used thing we deliver.

Roughly 18,000 to 26,000 dollars, fixed. That buys four to seven weeks of one person and most of it is not software. The audit, the interviews at each property and the editorial passes where your people decide the one right answer are the bulk of it.

Be careful with our own numbers. At our conservative assumptions this never pays for itself on prevented interruptions, which is why we print a threshold you can test rather than a saving we made up. We have no headcount for you, no revenue and no call volume.

The ask is one reply carrying one number, and it is at the end of this document.

01

This is ordinary, it is cheap to settle, and it is the proof

One defect you can check from outside, sixteen days before it matters.

Your Wedgewood Resort aurora page says the season runs 21 August through 21 April. Your Bear Lodge rooms page, describing the Northern Lights wakeup call and the headlamps you rent at the front desk, dates the same season mid-August to late April. The start dates disagree by about a week, though mid-August is a phrase rather than a date, so the size of that gap is our reading of your words.

This is not a sign of anything wrong with how you run hotels. It is what happens to any business that has maintained a lot of pages over a lot of years. It matters because of which window it is, and because settling it costs one decision.

21 August

Sixteen days from the date of this study, when the answer changes on two of your three properties. We found no aurora content published for Sophie Station Suites, which is itself worth knowing.

Your other two seasonal windows are tidy. The late check-out rule and the Golden Bear Restaurant are both published as mid-May to mid-September, the same span on both pages. The aurora window is the odd one out, it is the one your two pages disagree on, and it is the one that turns over first.

02

The cost is an interruption nobody counts, and it is time rather than money

It never appears on a schedule, so it never appears in a budget.

When a new clerk can't answer something, the question walks across the lobby to the most experienced person on the floor. That person is on the phone, or in the back, or is a seasonal hire themselves that week. Nothing about this shows up as a line item, which is exactly why it survives year after year.

Be careful how you value that. Your desks are staffed around the clock either way, so twenty minutes handed back to a clerk is capacity, not cash. It only becomes money if it defers a seasonal hire or avoids a backfill. If no hire is deferred, read the dollar column as zero and the benefit as better guest service.

We also can't size the problem from outside and we refuse to guess. The number circulating in this industry, that 40 percent of front desk calls go unanswered, is a claim by a vendor executive with no study behind it. It appears nowhere in our numbers.

Four measures, no baselines

Every number this study would want about your desks is currently uncollected.

Measuring the before is the first two weeks of the work.

Industry context, not a measurement of you. A survey of 282 hoteliers by the American Hotel and Lodging Association with Hireology over December 2024 and January 2025 found 65 percent reporting staffing shortages, with front desk the second most acute at 26 percent.

03

What changes is that the answer stops depending on who is standing there

One shift, walked step by step, before and after.

The outcome is stated as a change in behaviour rather than as a tool. A first-week seasonal hire gives the same correct answer a ten-year veteran would, at any property and in either season, without escalating and without a guest ever talking to a machine.

A shift on 21 August, the morning aurora season opens at Wedgewood and Bear Lodge

TODAY	AFTER
1 Clerk is asked whether the Northern Lights wakeup call is running yet. Two of your pages give different start dates.	The card for the wakeup call appeared on the page this morning, carrying one agreed date range that your team settled before the season opened.
2 Clerk checks whether late check-out is available now that summer feels over. The rule actually runs to mid-September.	The late check-out card is still showing, with its own end date printed on it, so the answer is right for another three weeks.
3 Guest asks whether the Golden Bear is open. The window is published on the Bear Lodge page, and we have no way to know what a clerk has in front of them mid-shift.	The restaurant card carries its own window and disappears on its own date, so the calendar does the remembering.
4 Guest asks something genuinely unusual. Clerk finds a supervisor.	Clerk finds a supervisor. The card says who to ask and your desks are staffed around the clock anyway.
5 A policy changes. We could find no dated, owned answer published anywhere a clerk could reach.	One named person changes one card. All three desks are saying the same thing that afternoon, and the change is dated and signed.

Row four is deliberately unchanged. Asking a person is the right answer for anything genuinely new, and a build that tried to replace it would be selling you something worse than what you have.

The row we did not change is the one worth reading twice. Your desks are staffed 24 hours by design, so asking a person is fast and it is the right thing to do at a hotel. Nothing we could build beats turning your head. What this beats is the three cases where turning your head fails, and we would rather define the win that narrowly than claim a bigger one you would catch us on by October.


What we would build, and why there is no AI in it

Cards, a calendar, and a sheet of paper.

A card for each question, holding one agreed answer, the property it applies to, the dates it is true, where it came from, who owns it and when it was last checked. A calendar file holding your season dates. A page per property that shows only the cards that are true today, with a plain word filter for finding one. A printed sheet of the twenty questions asked most, per property, per season. A log of what people looked for and did not find.

The season switch is a date, so a rule handles it and every case can be tested. 21 August is past or it is not. Which property a clerk works at is a setting, not a guess. The van rule that excludes tour groups and large groups traveling together is written policy, so that is a rule too. Deciding the one right answer is editorial work your people do once, and no model should do it for you, because the answer has to be yours.

A card can go out of date, so it says so on its face. Past its review date it stays visible, because hiding the only answer a clerk has helps nobody, and it carries a band naming the owner and the date. It drops off the printed sheet automatically, so a stale answer never gets taped to a desk.



Your answers, dated and owned, become one page per property and one sheet of paper per shift. Anything that changes today stays where it already lives, because nothing here connects to it.

There is one thing this will never do, and it is structural rather than a limitation we plan to fix later. The cards hold policy and how-to, never anything that changes today. No availability, no rates, no room status. Nothing we verified names a booking or property system you run, so a card claiming to know tonight's rooms would be a card that lies. A real share of what your desk gets asked is exactly that kind of question, so the honest ceiling on this is lower than a pitch would imply.

Where the industry evidence points matters here. The categories with real measured results are call triage, guest messaging, pricing, housekeeping routing and review response. What is NOT proven is a model in front of a guest making decisions. In a study of 340 UK adults who had used chatbots to book hotels, published in the International Journal of Hospitality Management, the discomfort those chatbots caused cut willingness to keep interacting by nearly 38 percent, and inaccuracy about room rates and cancellation policies nearly doubled the chance a booking was delayed or abandoned. A vendor survey of more than 500 properties found 59 percent of hoteliers think the front desk welcome should stay human-led. Applying a guest-facing finding to a staff-facing tool is our inference and not a measured result, and we would rather flag that than borrow the number quietly.

Some of this you should buy rather than let us build, and it is worth knowing what is actually working at properties your size before you spend anything.

The closest published case to you is a 96 room family owned lodge in Stowe Vermont, remote enough that the average staff commute is an hour, which is why they started. They cut call volume 30 percent and took the average guest message reply from 9 minutes to under 30 seconds. A 104 key family co-owned resort next to Yosemite had a problem you would recognise, a large share of its roughly 250 daily calls were people asking about the park rather than guests booking rooms. Their owner's stated win was not fewer staff, it was no longer rostering full shifts to cover two peak hours.

On pricing, an independent three property coastal group running seasonal 44 to 49 room properties reports occupancy up 7.8 percent and revenue per room up 8.4 percent over a year, plus six to twelve hours a week back for the owner. We quote that number rather than the 19 percent the same vendor markets, because the marketed figure has no control group and the named case does.

Housekeeping route software is the one with the hardest currency behind it, one operator cut a property's turnaround from three hours to one hour forty eight, worth about 27,000 pounds a year at that single site. For your seasonal crew the value is

as much in supervisor time and consistency as in labour dollars. Buy it, do not let us build it.

Two honest qualifiers. Most of these numbers are published by the vendor about a named customer, which is real information and a selected sample, because nobody publishes the deployment that did nothing. The upsell figures you will see quoted also overstate the AI contribution, most of that lift comes from putting an offer in front of a guest before arrival at all, which is a workflow change. We could not find a single Alaska hotel with a published measured deployment, so the nearest analogs really are Vermont and Bass Lake.

05

The conservative case does not pay for itself on prevented interruptions alone

So we print the threshold instead of a saving, and you can test it in a week.

We can't project a labour saving here and we are not going to pretend to. Every input such a model needs is missing from anything we could verify. Nothing in the table below was measured at your desks, and the marks say so on every row.

Read the dollars as conditional. Prevented interruptions are capacity, and they turn into money only if the hours defer a seasonal hire or avoid a backfill. So the table answers one question. How many interruptions per shift would this have to prevent before it has paid for itself over five years.

Break-even threshold, five-year view. Conservative pairs the high end of our fee with the lowest benefit assumptions.

	Conservative	Most likely	Aggressive
📌 Interruptions prevented per shift	2	4	5
📌 Desk shifts covered per day	3	6	9
📌 Minutes an interruption costs	4	5	6
📌 Loaded hourly rate	28	32	36
📌 Five-year total cost	43,094	38,374	33,654
📌 Annual value if those hours defer a hire	4,087	23,358	59,130
📌 Share of five-year cost recovered	43%	289%	862%
📌 Break-even interruptions per shift	4.7	1.4	0.6

Read the last row first. At our conservative assumptions this would need to prevent 4.7 interruptions per shift to cover its five-year cost, and we only assumed 2, which is why that column recovers 43 percent and never pays back. Not one row is verified, because not one figure was measured at your desks, and the middle column is the midpoint of our assumed ranges rather than a probability judgment. Desk shifts means three eight-hour shifts at one property, so the conservative column assumes one property adopts and the other two watch, and the aggressive column assumes all three. Year one counts about six months of benefit in the conservative column, nine in the middle and eleven in the aggressive, because a slower start is part of what makes the conservative case conservative. Fee 18,000 to 26,000, about 40 hours of your people's time across the phase, and upkeep of about 48 hours a year at an assumed 38 dollars plus a small reprint and hosting allowance, which rounds to 2,000 a year and is counted from the middle of year one. Contingency 18 percent.

📌 modelled from stated drivers 📌 assumed, and named as such

43%

of five-year cost recovered in the conservative case, against 289 percent in the middle column. If a hire is deferred, payback never comes in the conservative column, arrives at roughly a year and a half in the middle and roughly six months at the top.

The honest starting assumption for any AI proposal, including ours, is that it will not pay. MIT's Project NANDA put roughly 95 percent of enterprise AI pilots at no measurable profit impact, and RAND found more than 80 percent of AI projects fail. Neither is a fact about you. The strongest argument this build beats them is also the least flattering to us. There is no AI in it, so most of the documented failure modes structurally can't apply. The base rate you should actually fear here is the ordinary fate of internal documentation, beautiful in month one and wrong by month eight, and only a named owner and a review tied to dates you already publish touch that risk.

Your Operations Manager should own the number, with one front-desk supervisor owning the cards day to day. We name the role and not a person, because the only source we have for who holds it is a profile from 2018. If no single name can be attached to the upkeep hours, that is a reason not to start rather than a detail to settle later.

06

How we would start, and where you can stop

Gates all the way down, each one your off-ramp, and you keep something real at every one.

You can stop at any gate. What follows assumes the pre-quote question has already been answered.

Phase	What runs	How we would know it worked
Now	Ask the returning-versus-new question, before any quote	One number, the returning share of the summer roster
	Two weeks of paper tally at three desks, before anything is built	Escalations per shift, and the share of questions that are live status
	Audit every answer you publish, across three properties and both seasons	Count of contradictions found, and one agreed aurora date range before 21 August
	Turn what you publish into cards with property, dates, source, owner and review date	Cards live with an owner, and zero disagreements with your own website
	Write down the operating long tail that is on no page	Share of a shift's real questions a card now answers
	The printed shift sheet, summer and aurora, taped to each desk	The aurora cards appearing on the Wedgewood and Bear Lodge sheets on 21 August without a reminder
	Hand it over with a named owner and a plain export	One review run by them with us watching, not doing
Next	A search box that matches a guest's wording to your card, priced only if your own log asks for it	A repeating set of searches where the right card existed and the wording missed it
	Written answers for the multi-part questions that actually recur	Escalations on those patterns, before and after

You can stop at any gate and keep what it produced. After the audit you keep the full list of contradictions and one settled aurora date range, and you can fix your own pages and never speak to us again. After the cards exist you keep the whole set permanently, no licence, nothing that switches off. After the season run you keep four baselines and the question log.

Phase	What runs	How we would know it worked
Later	Buy housekeeping route software, pilot one property before three	Minutes saved per room clean on your property, at your occupancy and your labour rate
	The extended-stay and corporate reply path	Enquiry volume and current first-response time, which only you can measure
	A form and a shared board for courtesy van arrivals, not a voice agent	Share of arrivals reaching the desk as a complete record
	Sorting the general inbox by property and intent	Time from a message landing to reaching the right property

You can stop at any gate and keep what it produced. After the audit you keep the full list of contradictions and one settled aurora date range, and you can fix your own pages and never speak to us again. After the cards exist you keep the whole set permanently, no licence, nothing that switches off. After the season run you keep four baselines and the question log.

A named owner with authority to settle a disputed answer. About two hours each from a veteran at each property. Whatever internal training material already exists, which is the biggest reason the fee moves within its range.

07

What would make us wrong

Four ways this is a bad buy, and we would rather you found them now.

Your summer front desk might be mostly the same people every year, which is why we ask before quoting rather than after. It would shrink the prize sharply without zeroing it, because returning staff still come back after eight months away, across the switch

between your two seasons, and the date difference on your own pages is a defect regardless of who reads it.

The thing that actually gets used might be the sheet of paper, not the software. We think this is the likely outcome and we designed for it, which is why the printed sheet is a real deliverable rather than an export nobody maintains. If the screen goes unused your log will show it, and that is the strongest possible argument against funding anything smarter on top.

Most of what your desk gets asked might be live status, which a card set can never answer, and your own two-week tally will tell you what share before you commit to anything.

Nobody might own it. An answer book with no owner goes stale, and a confidently wrong card taped to a desk is worse than the honest gap you have today. If no single name can be attached to about 48 hours a year, tell us at the first gate and do not start. There is also a version that costs you nothing. Run the tally yourselves for a week and fix the aurora date on your own pages. If that is enough, keep your money.

08

The next step

One reply telling us roughly what share of your summer front-of-house people worked a season with you before. That single number decides whether any of this is worth your time, and we would rather you spent it on the answer than on a call.

What we checked

Every number in this study traces to a page we read. If we have any of it wrong, that is worth knowing too.

- 1** Your jobs page, on doubling staffing for summer and the seasonal roster
<https://www.fountainheadhotels.com/jobs>
- 2** Your FAQ, on check-in, late check-outs, cancellations, 24/7 desks and the shuttle call
<https://www.fountainheadhotels.com/faq>
- 3** Your Wedgewood aurora page, dating the season 21 August through 21 April
<https://www.fountainheadhotels.com/wedgewood-resort/specials/aurora-season-special>
- 4** Your Bear Lodge page, on a seasonal hotel and the Golden Bear window
<https://www.fountainheadhotels.com/bear-lodge>
- 5** Your Bear Lodge rooms page, on the courtesy van, its group exclusion, and the wakeup call dated mid-August to late April
<https://www.fountainheadhotels.com/bear-lodge-rooms.html>
- 6** Your contact page, listing all three properties
<https://www.fountainheadhotels.com/contact-us>
- 7** Your homepage, on extended stay routed to a named reservations manager
<https://www.fountainheadhotels.com>
- 8** Explore Fairbanks listing for Wedgewood, the third-party source we decline to rescale from
<https://www.explorefairbanks.com/listing/wedgewood-resort-fountainhead-hotels/10108/>
- 9** Study of 340 hotel chatbot users, International Journal of Hospitality Management
<https://techxplore.com/news/2026-05-hotel-chatbots-customers.html>
- 10** AHLA and Hireology survey of 282 hoteliers, fielded December 2024 to January 2025
<https://www.ahla.com/news/65-surveyed-hotels-report-staffing-shortages>
- 11** Independent eight-week test of housekeeping route software

<https://6hoteliers.com/housekeeping/optii-review/>

- 12** The 40 percent unanswered-calls figure, stated by a vendor executive with no study
<https://www.hoteldive.com/news/hotel-tech-ai-voice-front-desk-calls/740672/>
- 13** MIT Project NANDA, State of AI in Business 2025, on enterprise AI pilots showing no measurable profit impact
<https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/>
- 14** RAND, The Root Causes of Failure for AI Projects
https://www.rand.org/pubs/research_reports/RRA2680-1.html
- 15** A 96-room family-owned Vermont lodge, call volume down 30 percent and guest reply time from 9 minutes to under 30 seconds
<https://www.canarytechnologies.com/customers/trapp-family-lodge>
- 16** A 104-key family co-owned resort beside Yosemite, on absorbing non-guest call volume
<https://www.canarytechnologies.com/customers/the-pines-resort>
- 17** An independent seasonal three-property coastal group, occupancy up 7.79 percent and RevPAR up 8.40 percent over a year
<https://roompricegenie.com/case-studies/heron-hospitality/>
- 18** Housekeeping routing, one property's turnaround from three hours to one hour forty-eight
<https://www.optiisolutions.com/case-studies/how-staycity-saved-one-hour-fortyfive-minutes>