

Your site already tells patients when the phones are busiest, and that sentence is the most useful thing we found.

ANCHORAGE, ALASKA AUGUST 1, 2026

IN SHORT

Medical Park publishes something most practices keep private. Your frequently asked questions page says "Our peak calling time is between 8:30am-10am every morning, especially Mondays." You then tell patients to use the portal or call after 11. That is a practice that knows its own bottleneck and is honest with people about it.

We spent this week on what travels down that line. A good deal of it has one fixed answer that you have already written down. A refill starts at the pharmacy, which contacts you, and you ask for 72 hours. A records request takes seven to ten days, free if a physician asks or if it is just the last visit, \$25 for a full chart. Ten insurance plans are accepted and five are closed to new patients. Saturday is walk-in only, for acute illness and injury. Records before October 1, 2019 sit behind a separate login. Every one of those answers exists in your words on your pages. Reaching them today means reaching a person, during the hour you named as your busiest.

What we would build first has no AI in it at all. A pre-call check on your own site that answers the fixed questions, and a structured callback queue for the ones it can't, so your front desk works a ranked list instead of whatever rings next. Two weeks,

fixed price, roughly 13,500 dollars. It is plain rules your staff can read, it can be tested exhaustively, and it can't invent an answer.

Here is the honest counterweight, and we would rather you hear it from us. Your frequently asked questions page already publishes most of those answers, and the morning peak persists anyway. That is real evidence against the idea that a tidier version of your own content will move your phones. It is exactly why the ask is two weeks and measured rather than two months and confident.

What we could not verify is most of it. No call volume, no hold times, no registration counts are published anywhere we could find, so every number in our cost table is one we chose and named, not one of yours. Week one of the work is counting, and the count becomes the number both sides argue from.

We are also telling you not to buy two things. Not an ambient scribe, because athenahealth's chief executive has said it goes into your core athenaOne fee at no incremental charge with testing this year. Not anything that answers your phone, because the arithmetic on that is bad and we show it below.

The ask is a reply. No call and no pitch.

01

You already tell patients when the phones are busiest, which is rarer than it sounds

Most companies make you go looking for the operational gap. Medical Park printed it, then printed the workaround.

The sentence is on your frequently asked questions page. "Our peak calling time is between 8:30am-10am every morning, especially Mondays." The next sentence points patients at the portal or asks them to call after 11. Read that as what it is, a practice paying enough attention to its own front desk to warn people about it in writing.

Almost everything has to travel down that line, much of it by design. Your contact page tells patients that for clinical questions or appointment requests they should use the phone or the Patient Portal "rather than email" because "Email communication is not secure or HIPAA compliant." That is correct, and we want to be clear we think so. Your portal, in your own description, lets a patient request and confirm an appointment, which is a request a person then works. Saturday is walk-in only, so Saturday has no appointment channel at all.

8:30 to 10am

The window your own page names, especially Mondays. We could not find a published number for how many calls arrive in it, and we did not invent one.

Some of what arrives in that hour has exactly one correct answer, already written on your site. How a refill really moves, starting at the pharmacy, with 72 hours asked. What a records request costs and takes, seven to ten days, with no charge if a physician requests them or if it is only the last visit, and \$25 otherwise for a full chart release. Which of ten plans you accept and which five are closed to new patients. Where records before October 1, 2019 live. None of those needs a person to look anything up. All of them need a person today.

02

We can describe the shape of this, and we can't price it, because you have not published the numbers

This is the part where a vendor would show you a chart. We are not going to, because the inputs would be someone else's.

Nationally, staff phone time in medical practices splits into eligibility and prior authorization at 45 percent, scheduling at 31, intake at 9, prescription refills at 6, and an other category at 9. That is a real figure from a poll of 294 medical group leaders. It is

also not yours, and using it as your baseline would be the single most common dishonest move in this business.

So take the shape instead. Five roles are open on your jobs page against a team your about page puts at over 90 employees, and the Registered Nurse and Radiologic Technologist postings have been up since late April, roughly three months. When hiring is the binding constraint rather than budget, an hour spent repeating an answer that is already published is an hour you can't buy back at any price.

One of those five postings suggests somebody there is already asking this question. Your Executive Assistant role, published July 7, asks for "Process improvement experience and use of AI tools preferred".

Five open roles

Two of them, Registered Nurse and Radiologic Technologist, posted April 30 and April 29 and still listed on August 1.

That hiring pattern also changes how we are allowed to talk about return. Time given back to your front desk is capacity, not cash, because the constraint is finding people rather than affording them. We label it that way throughout, and we will not relabel it unless you tell us this would let you stop recruiting for a role.

03

What changes is that a fixed answer stops needing a person at 8:40 on a Monday

Walk one real call through, as it runs today and as it would run after. The step we chose not to change is the interesting one.

The outcome we would aim at is deliberately narrow. Fewer calls in your published peak whose answer was already on your site, and a structured place for the calls that genuinely need a person to land. Nothing clinical, nothing booked, nothing after hours.

One patient asking what a full chart release costs, Monday at 8:40

TODAY	AFTER
1 Patient calls the main line during the hour your page names as peak.	Patient opens the pre-call check on their phone.
2 Call joins the queue behind same day requests and refill questions.	Picks records request from a short list of what people actually call about.
3 Front desk answers, listens, and looks up the fee and the timeline.	Reads your answer, seven to ten days and \$25 for a full chart release, in your words.
4 Staff member explains that records before October 2019 sit in the older system.	Same answer, on the page, with the date it was last reviewed shown.
5 Patient asks a follow-up the page had not anticipated.	Patient asks a follow-up the page had not anticipated.
6 Staff member takes the question and promises to call back.	Question lands as a structured callback with a reason, a number and a best time.

The Today column is our reconstruction from what you publish, not from watching your desk, and week one of the work replaces it with what actually happens. Row five is identical on purpose. A patient with a question nobody anticipated should reach a person, and we are not going to put software between them and your staff to make a figure look tidier.


The callback queue matters more than the page. A page that answers questions with nowhere to put the ones it can't just moves the problem to a different desk. The queue is where deflected work lands as a filled-in row rather than a voicemail, and it produces the first structured data anyone here has about what those mornings are made of.

Most of what we would build has no AI in it, and the first phase has none at all

A page, a list of answers you own and edit, a queue, and a counter. That is the whole of phase one.

A short page on your own site that answers the questions with one fixed answer. A file of those answers that your staff edit without calling us, each with a review date. A queue where anything the page can't answer lands with a name, a number and a reason. A counter that records what was asked. Later, and only maybe, a reader that pulls text off the paper forms people still hand you.

A patient opens the check on a phone, picks what they need, and gets the answer you already publish. The answers live in a versioned file you own rather than inside code, so correcting a line is an edit, never a deploy. When the page can't answer, it collects the question into a callback queue your front desk works from one ordered list. Every interaction writes a row to a counter, which is how you get a number for what your 8:30 to 10am calls are actually about. Any clinical question, any symptom, anything the page is unsure of, stops and routes the person to your phone or your portal. It never guesses and it never reassures.



Read left to right. Phase one is the top path and it has no model in it. Anything the page can't answer becomes a structured callback your front desk works from a list. The paper form reader is the only place a model appears, it belongs to phase two, and it proposes to a person rather than writing anything. Notice that the only line into your record system starts at your front desk, not at the reader.

We would not put AI on your phones, in your after-hours line, on your prior authorizations, or anywhere near your coding, and we would not sell you an ambient scribe at all.

The phones are the clearest case. Answering a call end to end is roughly ten judgement steps, and chaining judgement steps multiplies their error. At 95 percent per step that is about 60 percent of calls completing correctly, at 90 percent about 35. Those per-step figures are ours, not measured on your line, and we print them so you can substitute your own. What is not in doubt is the direction, because the arithmetic compounds against you on the one channel where a missed acute complaint is a safety event rather than an inconvenience.

Prior authorization is the largest slice of that phone time at 45 percent nationally, and it is the biggest thing we could have proposed to you. We are not proposing it, because not one page of yours mentions it and we will not use a national average as your number.

Your after-hours arrangement, a staffed pediatric triage line and a named human backstop for adults, is more conservative than what most vendors would sell you, and we think you should keep it exactly as it is.

On the scribe, athenahealth's chief executive has said ambient documentation "will all be embedded in the core athenaOne fee and will not be at any incremental charge" with testing beginning in the first half of 2026. We are not telling you those tools do not work. A randomised trial found one cut time in note by 9.5 percent, from 4m30s to 3m49s, while another showed no significant change. We are telling you not to buy it from anyone, us included, when your own vendor has said it is about to include it. The call to make is to your athenahealth representative.

We buy the hosting, the certificates and the backups, because those are commodities. We build the page, the answer file, the queue and the counter, because they are small and you should own them.

If a second phase ever happens, we would ask athenahealth what forms capability your contract already includes before we write a renderer, for the same reason we told you to wait on the scribe. If they already cover it, our work shrinks to your own paper templates, the checkbox reading and the connection into the chart. We would rather lose that scope than sell you something you already pay for.

The conservative case clears, and it clears mostly on avoided mistakes rather than saved minutes

Three scenarios. Every figure below is computed from the drivers printed underneath it, and not one row is a fact about your business yet.

Read the marks in the left column before the numbers. Nothing here is marked verified, and that is the honest state of the evidence rather than an oversight. No call volume, hold time or registration count is published anywhere we could find, and we refused to put a national average in a column with your name on it. Week one replaces the most important assumption in this table with your own count.

Five-year total cost of ownership against modelled benefit, for the two-week phase one ask only. Phase two is not costed here and must not be read into it.

	Conservative	Most likely	Aggressive
ⓘ Front desk hours returned per working day	0.6	1.4	2.6
ⓘ Fully loaded front desk hourly cost	\$32	\$34	\$36
ⓘ Repeat inbound calls removed per year	1,500	4,200	8,000
ⓘ Chance the write-back path does not exist	20%	30%	35%
ⓘ Cost avoided if it does not, at that chance	\$44,000	\$52,000	\$60,000
ⓘ Chance of one wrong purchase made without a baseline	15%	25%	35%
ⓘ Cost of that purchase, at that chance	\$25,000	\$32,000	\$40,000
☑ Annual run rate benefit	\$19,750	\$43,480	\$76,000
☑ Five-year benefit, year one ramped	\$84,925	\$191,312	\$342,000
☑ Five-year total cost of ownership	\$25,110	\$24,692	\$24,064
☑ Share of five-year cost recovered	338%	775%	1,421%
☑ Payback month	21	13	6

How to reproduce every figure. Working days are 250 a year. Phase one is \$13,500 fixed, which is 2.5 person-weeks at a blended rate of about \$5,400 per person-week. That rate is our assumption and you should argue with it. Training is \$900 of our time to run it, and it also asks about six to seven hours of yours, which are separate quantities. Run cost is \$1,450 a year for hosting, monitoring and about two hours a quarter of content maintenance, and the maintenance is in there on purpose because your price list page says it was last updated January 30, 2024. Contingency is 20, 18 and 15 percent. Year one benefit is ramped to 30, 40 and 50 percent, so there is no day-one full benefit anywhere. The two avoided-cost lines are probability weighted, and here is the multiplication. In the conservative column that is 20 percent of \$44,000 and 15 percent of \$25,000, so about \$12,550 of the \$19,750 annual run rate. Most of the conservative case is avoided mistakes, not saved minutes, and you should read it that way. The hours per day figure is ours, not yours. At about 2.5 minutes of handling per routine call, the conservative column is roughly 14 deflected calls a day, which is deliberately low for a practice of over 90 employees. The hourly cost is anchored between your own published starting pay for a new-graduate CMA at \$22 and a new-graduate RN at \$37. The aggressive column is shown so you can see how wide the range really is, and it is not a case we would plan against.

 modelled from stated drivers  assumed, and named as such

Month 21

Conservative payback month, and the conservative case clears mostly on avoided mistakes rather than saved minutes. Most likely month 13, aggressive month 6.

Two things here cut against us. Strip the two avoided-cost rows and the conservative column is close to break even on deflected minutes alone, so it clears mostly because avoiding a wrong commitment is worth more than the minutes are. Payback also looks fast in every column because the ask is small, which is evidence of a small ask rather than a strong return.

On the outside view, roughly 95 percent of enterprise generative-AI pilots produce no measurable profit-and-loss impact, and in this sector 46 percent of medical group leaders report AI made providers more productive while 27 percent report no gains at all. The barriers they named run from resistance to change and inconsistent use through to tool accuracy and cost. This sits outside that reference class for one structural reason above all others. There is no AI in what we are asking you to buy.

Whoever runs the clinic day to day on the administrative side owns these numbers, not a clinician. Your site names your providers and your HR manager, and it publishes no administrator or executive title, so we are not going to assign one to anybody. They would track four things against the week-one baseline, calls in the 8:30 to 10am block, the share answerable from already-published information, queue items closed without a second inbound call, and minutes per new patient registration, then report actual against baseline in writing at ninety days, including the case where the page saw little use.

06

A free test first, then two weeks, then a decision you can walk away from

Sequenced by what is cheapest to learn from, not by what is
biggest to build.

Before the two weeks are committed, there is a free test. Your front desk marks a tally sheet for one week, including a Monday, and we agree the kill number before it runs. A test we would refuse to lose is not a test. Everything after that is judged against your own count rather than against a forecast we wrote.

Phase	What runs	How we would know it worked
Now	Run a five day call tally before you commit anything. One sheet at each phone position, one week including a Monday, nine buckets, plus a mark for could a page have answered this. Zero dollars, no software, and we write the kill number down first. Under 20 percent answerable and we tell you in writing not to build it.	One week of marks, alongside whatever call report your phone vendor already produces.
	The pre-call check, live on your site, answering only questions with one fixed answer, built from your own pages.	Completions per week and the change in published-answer calls against the tally. Target set after the baseline, never before.
	The structured callback queue behind it, so a request lands as a filled-in row rather than a voicemail.	Share returned the same business day, and median hours from request to callback. Baseline is zero because the channel does not exist.
	A five day investigation of the athenahealth write-back path and its business associate terms.	Three yes or no answers in writing, with the contract language attached.
	A named owner and a review date on every published answer, before the page ships.	A person named and the next review date already in a calendar.
Next	A paper to athenahealth intake pipeline for the forms patients actually fill in, with one AI step and	Minutes per registration against the baseline with review time counted

Three gates and each is a real way out. Gate zero, before anything is signed, the investigation returns a written yes or no on whether athenahealth permits a third party to write structured intake data, and on the business associate terms. If no, the honest recommendation becomes the page and the queue and nothing more from us. Gate one, the end of two weeks, is a complete stopping point. You keep the live page and its content, the callback queue, your measured baselines, the written answer on write-back, and a named owner with a review date. It is code and content you can read, host and edit, with no model in it and no dependency on us. Gate two, ninety days, judged against your own baseline and against a success bar you write down during week one, before the page goes live, so nobody can move it once the result is known.

Phase	What runs	How we would know it worked
	a person accepting every field.	in, never out, and the reading step measured field by field against a set of your own completed forms before it is trusted anywhere.
	Ask athenahealth what forms capability you already pay for, before we write a renderer.	A written comparison, which may well say buy rather than build.
	Count your prior authorization load for four weeks. You run this yourself, free, whether or not you hire us.	Authorizations started per week by payer and service line, plus staff minutes.
Later	Online self-scheduling and Saturday capacity visibility. Not from us. Ask athenahealth what your agreement already includes.	Share of routine bookings completed without a staff member.
	You already refuse email for clinical questions because it is not secure. A compliant channel that is not the phone is the gap that leaves, and it is a vendor and compliance decision before it is a technology one.	Share of routine questions arriving outside the phone.
	Anything that answers your phone, behind four weeks of measured volume, abandonment and reason mix.	Four weeks of call data, then a decision.

Three gates and each is a real way out. Gate zero, before anything is signed, the investigation returns a written yes or no on whether athenahealth permits a third party to write structured intake data, and on the business associate terms. If no, the honest recommendation becomes the page and the queue and nothing more from us. Gate one, the end of two weeks, is a complete stopping point. You keep the live page and its content, the callback queue, your measured baselines, the written answer on write-back, and a named owner with a review date. It is code and content you can read, host and edit, with no model in it and no dependency on us. Gate two, ninety days, judged against your own baseline and against a success bar you write down during week one, before the page goes live, so nobody can move it once the result is known.

Phase	What runs	How we would know it worked
	Ambient documentation. Wait for athenahealth and spend nothing. We will help you evaluate it when it lands.	Time in note before switch-on and again at ninety days, with the bar written down first.

Three gates and each is a real way out. Gate zero, before anything is signed, the investigation returns a written yes or no on whether athenahealth permits a third party to write structured intake data, and on the business associate terms. If no, the honest recommendation becomes the page and the queue and nothing more from us. Gate one, the end of two weeks, is a complete stopping point. You keep the live page and its content, the callback queue, your measured baselines, the written answer on write-back, and a named owner with a review date. It is code and content you can read, host and edit, with no model in it and no dependency on us. Gate two, ninety days, judged against your own baseline and against a success bar you write down during week one, before the page goes live, so nobody can move it once the result is known.

One named person who owns the answers and can be reached inside a day. About two hours from whoever answers the phone between 8:30 and 10am, plus the week of tally marks, which is the dependency everything rests on. A yes or no on whether any person keys a hand-filled form into athenahealth today. Your athenahealth account contact and permission to ask them directly. A check that each answer we copy is still current, particularly the price list.

07

Four things that would make us wrong, and the free test we would run first

We would rather raise these than have you find them.

It may not move your phones at all, and this is the assumption everything here rests on. We are assuming a meaningful share of your 8:30 to 10am peak is people asking things whose answers are already fixed and already published. The evidence cuts against us before we start. Your frequently asked questions page already carries most of those

answers and the peak persists anyway. Nationally the largest slice of practice phone time is eligibility and prior authorization at 45 percent and scheduling at 31, and neither is answerable by a page, because is my plan active and can you fit me in Thursday are both specific to one person. That is why we would run the tally first, for nothing, with the kill number agreed in advance.

Every number in our table is one we chose, because nothing about your volumes is published anywhere we could find. The conservative case also leans on avoided mistakes rather than saved minutes, so strip the two avoided-cost rows and it is close to break even on deflected time alone. If your real count comes back below that column, the honest answer is that this was not worth doing, and we would say so.

A second phase may be impossible, because whether athenahealth lets a third party write into your chart is a commercial question about a contract we have never seen. There is also a question that could delete the largest part of any second phase, so we will ask it before you do. Does any person here key a hand-filled form into athenahealth today? We guessed that somebody does, from the forms on your printable documents page that a patient fills in by hand. That is a guess about your back office, not a fact, and you can settle it in one sentence. If nobody keys them, the second phase is off.

08

The next step is small

Read it and tell us if it landed. We are not asking for a meeting. If you want the tally sheet or the four week prior authorization count, we will send both whether or not you ever hire us.

What we checked

Every number in this study traces to a page we read. If we have any of it wrong, that is worth knowing too.

- 1** Peak calling time is 8:30am to 10am, especially Mondays, with the portal or calling after 11 suggested
<https://mpfcak.com/patient-information/frequently-asked-questions/>
- 2** For clinical questions or appointment requests, use the phone or the Patient Portal rather than email, because email is not secure or HIPAA compliant
<https://mpfcak.com/contact-us/>
- 3** Records requests take seven to ten days, and a full chart release is charged at \$25
<https://mpfcak.com/patient-information/frequently-asked-questions/>
- 4** A refill starts at the pharmacy, which contacts the practice for authorization, with 72 hours asked
<https://mpfcak.com/patient-information/prescription-refills/>
- 5** Records before October 1, 2019 sit in a separate record system behind its own login
<https://mpfcak.com/patient-information/patient-portal/>
- 6** The patient portal is athenahealth, and a second athenahealth login holds records before October 1, 2019
<https://mpfcak.com/patient-information/patient-portal/>
- 7** The portal lets a patient request and confirm appointments, review results and pay a bill
<https://mpfcak.com/patient-information/patient-portal/>
- 8** Ten insurance plans accepted, five categories closed to new patients
<https://mpfcak.com/patient-information/insurance/>
- 9** Saturday is a walk-in clinic only for acute illnesses and injuries, and same-day slots are reserved with a request to call ahead
<https://mpfcak.com/patient-information/frequently-asked-questions/>
- 10** Pediatric after-hours goes to a 24-hour nurse triage line staffed with pediatric nurse specialists, adults to a next business day nurse callback or the

Providence line

<https://mpfcak.com/patient-information/frequently-asked-questions/>

- 11** Over 90 employees, and local and physician-owned since 1971
<https://mpfcak.com/about-us/>
- 12** Five open roles, with Registered Nurse published April 30 2026 and Radiologic Technologist April 29 2026, and starting pay of \$37 for a new-graduate RN and \$22 for a new-graduate CMA
<https://www.mpfcak.com/jobs/>
- 13** Published procedure price list, stating it was last updated January 30, 2024
<https://mpfcak.com/patient-information/top-10-procedure-codes-pricing/>
- 14** athenahealth's chief executive on ambient documentation being embedded in the core athenaOne fee at no incremental charge, with user testing in the first half of 2026
<https://www.newsweek.com/athenahealth-throws-its-hat-in-ambient-scribe-ai-assistant-rings-access-health-10991353>
- 15** Randomised trial at UCLA Health, one ambient tool cut time in note 9.5 percent from 4m30s to 3m49s, the other showed no significant change
<https://pmc.ncbi.nlm.nih.gov/articles/PMC12265753/>
- 16** Where staff phone time goes: eligibility and prior authorization 45 percent, scheduling 31, intake 9, refills 6, other 9
<https://www.mgma.com/mgma-stat/phones-are-still-a-backlog-costing-medical-practices-time>
- 17** 46 percent of medical group leaders report AI made providers more productive, 27 percent report no gains, 14 percent unsure, 13 percent do not use AI
<https://www.mgma.com/mgma-stat/are-ai-tools-making-clinicians-more-productive-its-complicated>
- 18** Roughly 95 percent of enterprise generative-AI pilots produce no measurable profit-and-loss impact
<https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/>
- 19** Executive Assistant posting, published July 7 2026, asks for process improvement experience and use of AI tools preferred
<https://mpfcak.com/jobs/executive-assistant/>

- 20** Most recent news post March 20, 2025, and the one before it December 28, 2023
<https://mpfcak.com/patient-information/medical-park-news/>
- 21** Downloadable patient documents, a mix of forms a patient fills in and policy documents, several dated 2020 and 2021
<https://mpfcak.com/patient-information/printable-documents/>
-

Alaska AI. Prepared for Medical Park Family Care. Not for redistribution.