Alaska AI

FIELD STUDY PREPARED FOR R&M CONSULTANTS, INC.

Your construction administration posting names a documentation duty that protects your public funding.

ANCHORAGE, ALASKA   AUGUST 7TH 2026

IN SHORT

We read your public pages, your job postings, your news archive and your competitors, then spent the day on one question. Where could a small, honest build change how R&M works.

The answer we expected was proposals. Two senior seats in two different departments carry the identical duty, "Conduct research and write technical portions of proposals to assist in bringing in new work.", and a third carries "Developing proposals." That is the loudest signal in your public material and we are not selling you anything for it. You hired a Proposal Manager twenty four days before we started, and we could not find one named firm at any size attached to a measured proposal productivity number with a stated method.

What we found instead sits in a sentence from your Construction Administration posting: "Administer project documentation and audit trails to ensure funding participation". At a firm where nearly 95 percent of work is for public clients, that is a sentence about money. We have never seen an R&M project record, so we have no

> evidence you are missing anything. What we can see is that the duty exists and that a miss is not clerical.
>
> So we would build the smallest useful thing. Before each pay estimate goes out, a check reads the documentation you already produce and reports the contract required items it did not find. It has no AI in it. Every requirement lives in a rule your own construction administration lead can read and argue with. It never says all clear, because a false clear is worse than shipping nothing.
>
> Twelve thousand dollars, eight weeks, two gates, and the first is four thousand and a real stop point.
>
> Your volumes, hours and rates are published nowhere, so every driver is our assumption and yours to correct. On the cautious version this recovers 119 percent of five year cost and pays back in month 47, just under four years, which is poor by any normal standard. The middle version is month 15. If one person both prepares and reviews, the cautious column drops to 99 percent and never pays back. That benefit is about 93 hours a year, roughly two and a half work weeks, and it is capacity rather than cash. Two numbers turn most of this into your own figures, roughly how many reconciliations you run a year and roughly how long one takes.

01

One sentence in your own job posting describes work that protects money, not paperwork

Your Construction Administration posting names the duty.
The client mix is what makes it expensive.

Your posting for a Project Civil Engineer in Construction Administration lists, among the duties, "Administer project documentation and audit trails to ensure funding participation". Read that against your Who We Are page, "Nearly 95% of our work is for public clients", and it becomes a sentence about money. An item that never reached the record is a reimbursement conversation rather than a filing problem.

We have never seen an R&M project record, so we have no evidence you are missing anything at all. What we can see is that the duty exists, that a person is named as responsible for it, and that the consequence of a miss is money.

The same posting asks for work "in project locations throughout the state for extended periods of time during construction season", and is classified Seasonal, Full-Time, Hourly. That is ordinary staffing for an Alaska season and reads to us as seasonality rather than a gap.

One of eleven roles open across your four offices sits in that group. Your postings carry no dates, and eleven openings at a firm ranked tenth in its size bracket on Zweig Group's 2026 list reads as scale rather than trouble.

> **Nearly 95%**
>
> of your work is for public clients, in your own words on your own page. That is the number that turns a documentation gap into a funding question.

We checked the field in both directions. DOWL published a 2020 piece on machine learning that claims no deployment of its own, and beyond that we found no current AI position and no named tool at DOWL, CRW or RESPEC. That absence covers only the pages we fetched, and a firm can run AI internally and say nothing about it. A fourth competitor's site would not load across three attempts, so we assessed three of four.

The more useful half is where this is measurably working, and it is not where the marketing points. The strongest independent result in your industry is a Louisiana state transportation department evaluation of photo based field inspection across more than fifty construction projects, which measured a 28 percent productivity increase for inspectors creating and submitting daily work reports. Two caveats belong in the same breath. That system was photographs filed on a fixed form and routed through set steps rather than a model, which is part of why what we would build has no model in it. It also measures field inspectors writing daily reports, the spring 2027 item rather than phase one, so none of that 28 percent is in our table.

HDR publishes an AI statement that is mostly restraint, including "Prohibit the use of generative AI for use in any final work product". Michael Baker announced a platform in April reaching "more than 6,000 engineers, architects and scientists". Those are the economics of a rollout to a firm many times the size of the 100 plus people your own page names, and they do not transfer down.

02

We can't tell you what this costs you, and neither can anyone else from outside

Four numbers decide this and you are the only source for any of them.

We went looking for what this costs R&M and came back with nothing. Not the reconciliations you run in a year, not how long one takes, not what an hour costs you loaded. Your closeout count and what an assembly costs are not published either.

Nothing published says R&M has ever had a funding disallowance, an audit finding or a record dispute. We looked, and we are not hinting there is one. The cautious column assumes zero.

> **17 assumptions**
>
> Four move the answer most, how many reconciliations you run, how long one takes, what an hour costs you loaded, and how much of that time a check actually removes. The rest sit in the table below, every one marked. Your own four week baseline, started in week one, replaces the first two with measurements you take, and your own accounting supplies the third. How much a check removes is only known once it runs, at week eight.

03

What changes is that the list of what is missing exists before anyone goes looking

One monthly pay estimate submittal, walked step by step.

One caution before the figure. We have never seen your records or your process.

Preparing one monthly pay estimate submittal, today and after phase one

	TODAY	AFTER
1	~~A progress payment comes due, so someone opens the project folder and works out what this contract requires for this pay period.~~	The check has already run against that project's rulebook. The required list exists before anyone opens the folder.
2	~~They search daily reports, materials certifications, test results and quantity records to see what is actually there.~~	The report names the specific items the rules did not find, each one pointing back at the requirement it came from.
3	A senior reviewer builds a second completeness check from scratch before it goes to the owner.	The senior reviewer reads a short list instead of rebuilding the check.
4	Whatever is missing gets chased, sometimes after the crew has already left that site.	Whatever is missing was flagged before the submittal, while people are still out there.
5	A person decides whether the gap actually matters and what to do about it.	A person decides whether the gap actually matters and what to do about it.

The left column is our assumption, not your process. Row five is identical on purpose, and it is the most important row here. Deciding whether a gap actually matters is a judgement about a contract and about money, and it stays with the person whose name is on the project. We are not automating that and we would not know how to.

The figure leaves out the part that makes this safe. The report never tells you a project is complete. It lists what it did not find, then names what it could not look at, so a folder it failed to reach shows up as unexamined rather than passing.

There is a third thing it can't show you. A requirement nobody wrote into the rulebook produces no line at all, which is why gate B runs the rules against submittals your own reviewer already checked by hand, and why the rulebook has a named owner rather than an author.

04

What we would build, and there is no AI anywhere in it

A rulebook, a reader, a checker, and a report that admits what it missed.

A rulebook, one plain file per contract, holding what that contract requires, each rule naming its spec section. A reader that opens your project folders and writes nothing back. A checker that runs the rulebook against what it found. A report listing what it did not find, then naming everything it could not examine. Your construction administration lead owns the rulebook.

Ahead of each pay estimate or progress submittal, and again at closeout, it runs against what you already produce and reports the missing items with the requirement each came from.

Nothing changes for your field crews and there is nothing for them to install. The change lands on two desks, whoever prepares the estimate and whoever reviews it. If those two do not change how they work the hours in the table do not move.

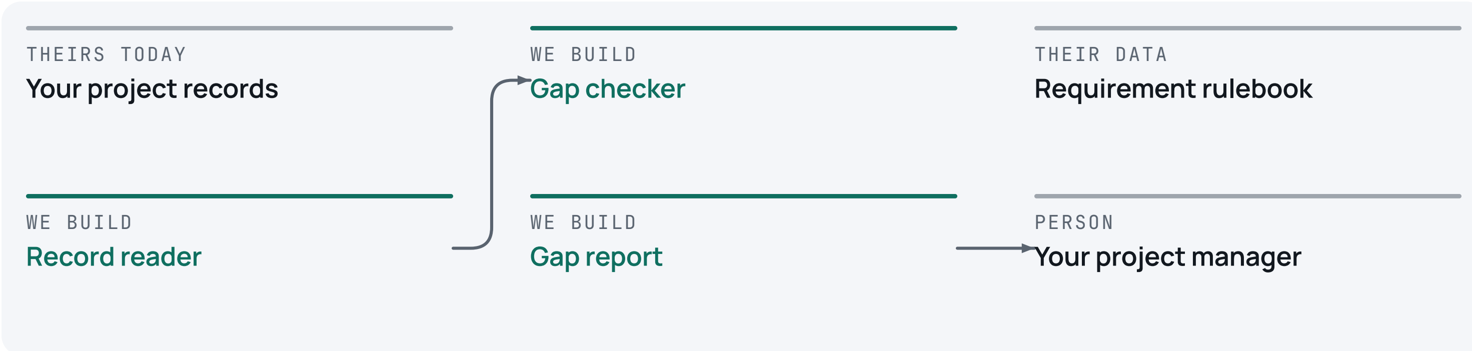

Phase one, end to end. The highlighted boxes are the work we would do. The owner agency's record system is deliberately not here, because phase one does not read it and gate A decides whether it ever could. There is no model box either. The arrow running back from your project manager to the rulebook matters most, since the rules are only as good as the person who argues with them.

A working demonstration, built on your own facts. It is a demonstration and it is honest about that, and it does nothing this study did not scope. Open it full screen.

An early draft put a small language model on one part of this. It is out. Documentation requirements are enumerable, so a rule does them cheaper and testably. A rule can be printed on a page and disputed by your construction administration lead. A model can only be sampled. Until somebody can name the requirement types the rules keep missing, a model here is unearned.

A firmwide assistant over 57 years of your records scored well on paper and is also out. Putting marked and unmarked federal material into one searchable pile runs against your CMMC Level 2 certification, announced in December 2025, which sets rules about where federal information sits.

One line holds at every phase. Nothing carrying a licensed professional's stamp gets reviewed by a model. ASCE Policy Statement 573 states that "The civil engineer must maintain responsibility for project planning, designing, building, operations, maintenance, and the protection of the public health, safety, and welfare". Insurers watch the same ground, with 80 percent of fifteen surveyed this year viewing AI adoption by design firms as a potential disruptor. Liability forms picked up generative AI exclusion endorsements in January 2026, optional rather than automatic, so that is a question for your broker rather than a settled fact.

We would buy every commodity part, the document parsing, the database, your own sign on, the report rendering, and hosting inside the environment your CMMC Level 2 certification already governs rather than standing anything new up beside it. That line is for whoever owns that boundary to draw rather than us. We would build only what is yours, starting with the reader, which has to understand folder conventions that will not be identical across Anchorage, Fairbanks, Juneau and Wasilla, then the rulebook format, the checker, and the piece that tracks what was not examined.

05

On the most cautious assumptions this pays back in month 47, just under four years

Nothing here is a fact about your business yet. That is what your own four week baseline is for.

Read the marks down the left before the numbers. Not one row is marked verified, because nothing about your volumes, hours or rates is published anywhere we could reach.

This ask started at eighteen thousand. At that price the cautious column never paid back, so the price came down rather than the assumptions going up.

A payback in month 47 is a poor result by any normal standard. Gate A costs four thousand and your own four week baseline starts in week one rather than after it, so your volumes and your hours get measured while the three questions are being answered. How much a check actually removes stays our assumption until it runs at week eight.

Five year view. Every input is a row, so the outputs at the bottom can be rebuilt from the page. Running cost starts at month seven, so it is charged for 4.5 of the five years, which is why the subtotal is 27,450 rather than 28,850. Your staff time is costed at the cautious hourly rate in all three columns, so it does not get cheaper as the case improves. Hours a year converts the closeout saving back to hours at the same loaded rate and leaves out the owner agency response, which we costed in dollars only, so the middle column is 121.5 plus 36 plus 70. Years three to five are held flat rather than compounded. Cost rows are ours and identical in all three columns.

	Conservative	Most likely	Aggressive
Our fee, phase one, fixed	$12,000	$12,000	$12,000
Your staff time, about 30 hours at the cautious rate, held flat across columns	$2,850	$2,850	$2,850
Running and maintaining it, $2,800 a year from month seven, so 4.5 years	$12,600	$12,600	$12,600
Subtotal before contingency	$27,450	$27,450	$27,450
Contingency on the whole stack	20%	20%	20%
Five year cost, all in	$32,940	$32,940	$32,940
Reconciliation events a year	60	90	110

Every driver is a row above and none came from you. The cautious column assumes zero owner agency record reviews, so nothing in it depends on you having had a problem. About a third of that column sits in the closeout rows, and a four week baseline may not contain a closeout, so those rows can still be our assumptions when gate C is decided. Recovered hours are capacity rather than cash. About 93 hours a year on the cautious column is roughly two and a half work weeks of one person's time, and only money if you redeploy it. Two things break the cautious column and we computed both. The two time recovery rows assume two different people. If one person does both, the column recovers 99 percent and never pays back inside five years. Separately, if running it overruns by a fifth, the column falls to 109 percent and payback slips to month 53.

	Conservative	Most likely	Aggressive
Hours per event, locating and cross checking	2.0	3.0	3.5
Loaded internal cost an hour	$95	$110	$120
Portion of that time removed	40%	45%	50%
Senior review pass, hours per event	0.75	1.0	1.0
Portion of the review pass removed	35%	40%	45%
Project closeouts a year	8	10	12
Documentation assembly cost per closeout	$1,140	$2,200	$2,800
Portion of closeout assembly removed	30%	35%	40%
Owner agency record reviews a year	0	2	3
Internal response cost per review	$4,000	$4,000	$5,000
Portion of that response removed	35%	35%	40%
Year one adoption ramp	45%	55%	65%
Hours a year this gives back	about 93	about 228	about 354

Every driver is a row above and none came from you. The cautious column assumes zero owner agency record reviews, so nothing in it depends on you having had a problem. About a third of that column sits in the closeout rows, and a four week baseline may not contain a closeout, so those rows can still be our assumptions when gate C is decided. Recovered hours are capacity rather than cash. About 93 hours a year on the cautious column is roughly two and a half work weeks of one person's time, and only money if you redeploy it. Two things break the cautious column and we computed both. The two time recovery rows assume two different people. If one person does both, the column recovers 99 percent and never pays back inside five years. Separately, if running it overruns by a fifth, the column falls to 109 percent and payback slips to month 53.

	Conservative	Most likely	Aggressive
Annual value once running	$8,792	$27,825	$48,480
Share of five year cost recovered	119%	384%	684%
Payback	month 47	month 15	month 7

Every driver is a row above and none came from you. The cautious column assumes zero owner agency record reviews, so nothing in it depends on you having had a problem. About a third of that column sits in the closeout rows, and a four week baseline may not contain a closeout, so those rows can still be our assumptions when gate C is decided. Recovered hours are capacity rather than cash. About 93 hours a year on the cautious column is roughly two and a half work weeks of one person's time, and only money if you redeploy it. Two things break the cautious column and we computed both. The two time recovery rows assume two different people. If one person does both, the column recovers 99 percent and never pays back inside five years. Separately, if running it overruns by a fifth, the column falls to 109 percent and payback slips to month 53.

modelled from stated drivers   assumed, and named as such

> **119%**
>
> of five year cost recovered on the most cautious assumptions, paying back in month 47, just under four years. That is about 93 hours a year, roughly two and a half work weeks, and it is capacity rather than cash until you redeploy it. The middle version is month 15.

The outside view here is bad and it applies to us. MIT's NANDA initiative reported in August 2025 that only about 5 percent of AI pilot programs reach rapid revenue acceleration, with the rest stalling at little or no measurable impact on profit and loss. That counts pilots without a bottom line result rather than AI not working, and we claim no exemption.

Four things here are attempts to be in the minority that pays. Phase one contains no AI, so its reliability is rule coverage, checkable against your own contract documents. It asks no adoption of your field crews. Your baseline is measured before the larger commitment, and the number gets a named owner and a check at 90 days.

Your construction administration lead should own the number, since it is their posting the whole finding came from, with whoever runs project accounting countersigning the hours baseline so the owner is not the only person who sees it.

06

You can stop after two weeks and keep everything it produced

Two gates inside the ask, and the first one is a genuine off ramp.

Phase one is twelve thousand dollars fixed, eight weeks, four thousand at gate A and eight thousand at gate B. The only irreversible commitment is the four thousand.

Gate A is two weeks and writes no code. It answers three questions we do not know. Whether an owner agency dictates the field record system on your federal aid work and whether you may feed it. Whether today's documentation is machine readable or sitting in scans. Which projects carry controlled markings. Two of those can go against us.

Phase	What runs	How we would know it worked
Now	Gate A, two weeks of questions and no code. Three answers in writing, sourced to your own contracts, plus a draft rulebook from one real contract	The three answers exist and are sourced. Success is that they are answered, not that they are favourable
	Your own four week baseline, started week one and owned by your construction administration lead. Hours per reconciliation and the count, plus the hours on the next closeout whenever it lands, since at eight to twelve a year one may not fall inside the four weeks	Four consecutive weeks logged by your people, roughly ten minutes per submittal. Without it, nothing later can be judged and we would say so
	Gate B, the rules only check running on two or three live projects through your 2026 closeout season	Zero misses on a run against submittals your own reviewer already checked by hand. Then the funding bar, which is that the four week baseline shows at least 2.0 hours per reconciliation event and the measured reduction reaches 40 percent, which is the cautious column
Next	The controlled marking and boundary decision, recorded before any component would cache project data on a field device	A written scope decision per component, and no new in scope system without a recorded change
	Field reports filled in on a form instead of typed free hand, at spring mobilisation	Hours per report day, and the share of reports submitted within 24 and 72 hours

Gate A, end of week two, is a real stop. You keep the answers and the started baseline, and the eight thousand is never invoiced. Gate B, end of week eight, is also a stop, and you keep a running check, the rule set your lead signed, and a before and after on your own numbers. Gate C is the winter decision, by March 1st 2027, and Alaska sets that date rather than us. The spring phase gets funded only if the four week baseline shows at least 2.0 hours per reconciliation event and the measured reduction reaches 40 percent, which is the cautious column. Below either, phase one is the finding and the roadmap stops there. That comes off two or three projects, a handful of events rather than a season, so a borderline result reads as directional rather than as the bar cleared.

Phase	What runs	How we would know it worked
	2027, while seasonal staff are onboarding anyway	
	An export into the owner agency format that produces the same file from the same records every time, built only if gate A found an export path exists	Share of records accepted on first submission with no rekeying
	One narrow model step, a coverage flag on requirement types the rules keep missing	A named list of missed requirement types from a real closeout cycle. A short list means it does not ship
Later	An evidence locator over one project's record for when a change order or claim arrives months later. It returns documents and citations, never conclusions	You state the frequency first, since we could find no published figure for how often that reaches you
	NEPA and permitting drafting, two of your own 24 service lines, capped at 1 to 5 hours per subsection, which is the vendor's own published benchmark rather than an independent measurement	Hours per subsection against your baseline, plus the rate at which stale references get caught
	Grant writing, a service line you sell rather than an overhead, which is the most interesting thing we did not pick	You supply grant volume and win rate first. Until then it stays a hypothesis

Gate A, end of week two, is a real stop. You keep the answers and the started baseline, and the eight thousand is never invoiced. Gate B, end of week eight, is also a stop, and you keep a running check, the rule set your lead signed, and a before and after on your own numbers. Gate C is the winter decision, by March 1st 2027, and Alaska sets that date rather than us. The spring phase gets funded only if the four week baseline shows at least 2.0 hours per reconciliation event and the measured reduction reaches 40 percent, which is the cautious column. Below either, phase one is the finding and the roadmap stops there. That comes off two or three projects, a handful of events rather than a season, so a borderline result reads as directional rather than as the bar cleared.

Phase	What runs	How we would know it worked
	Materials testing and special inspections need plain automation and not AI. Breaks, gradations and densities are structured and high volume	Share of results reaching the client report with no retyping

Gate A, end of week two, is a real stop. You keep the answers and the started baseline, and the eight thousand is never invoiced. Gate B, end of week eight, is also a stop, and you keep a running check, the rule set your lead signed, and a before and after on your own numbers. Gate C is the winter decision, by March 1st 2027, and Alaska sets that date rather than us. The spring phase gets funded only if the four week baseline shows at least 2.0 hours per reconciliation event and the measured reduction reaches 40 percent, which is the cautious column. Below either, phase one is the finding and the roadmap stops there. That comes off two or three projects, a handful of events rather than a season, so a borderline result reads as directional rather than as the bar cleared.

Two or three named active projects, and a line from contracts on whether they carry controlled markings. Six to eight hours of your construction administration lead in weeks three and four. Read access to where the records live, scans included. Ten minutes per submittal from your project managers for four weeks. One named person who owns the number at 90 days.

07

Five ways this is a bad buy, and the first one could end it at week two

Every one of these is answered here rather than left hanging.

An owner agency may own the record and refuse a feed. That is the riskiest assumption here and it is unanswered today, which is why gate A costs four thousand rather than twelve. If it goes against us we tell you at week two and this becomes the whole build rather than its first phase. Phase one still works there, since it reads what you produce.

Your documentation may not be machine readable. If daily reports and materials records are scans, getting the text off them becomes the project and our estimate is wrong by a lot. Gate A measures it against a threshold we write down before we look. If it fails you stop at week two with the memo and a baseline two weeks in rather than four.

Our numbers may be wrong in either direction. The four week baseline replaces our volume and hours guesses with your measurements, the reduction rows stay assumptions until gate B, and the bar is stated in advance: the four week baseline shows at least 2.0 hours per reconciliation event and the measured reduction reaches 40 percent, which is the cautious column.

It may underperform even then. If the measured reduction misses 40 percent, gate C does not fund the spring phase. You stop there and you keep the running check and the rule set your own lead signed. The running cost stays yours from that point and you can switch it off.

The rules can go stale. Standard specifications and agency manuals get revised, and a list nobody maintains reports the wrong gaps against the new ones. Naming who maintains it is a condition of finishing.

08

Two numbers from you turn most of this table into your own figures

Tell us roughly how many pay estimate and progress submittal reconciliations you run in a year, and roughly how long one takes. Those two move about seven tenths of the cautious column.

What we checked

Every number in this study traces to a page we read. If we have any of it wrong, that is worth knowing too.

1. Your Construction Administration posting, on administering documentation and audit trails to ensure funding participation, on statewide travel during construction season, and on the Seasonal Full-Time Hourly classification
https://www.rmconsult.com/job-posting/project-engineer-construction-administration/
2. Your Who We Are page, on nearly 95 percent public clients, 100+ employees and founded 1969
https://www.rmconsult.com/who-we-are/
3. Your open positions page, the eleven listed roles, and the absence of any posting dates; also your four offices, Anchorage, Fairbanks, Juneau and Wasilla
https://www.rmconsult.com/join-us/open-positions/
4. The Senior Land Surveyor posting, on researching and writing technical portions of proposals
https://www.rmconsult.com/job-posting/senior-land-surveyor/
5. The Senior Project Civil Engineer posting, carrying the same proposal duty in a second department
https://www.rmconsult.com/job-posting/senior-project-engineer/
6. The Senior Environmental Geologist or Engineer posting, on developing proposals
https://www.rmconsult.com/job-posting/senior-environmental-geologist-or-engineer/
7. Your July 14th 2026 announcement of a Proposal Manager into the Marketing Group
https://www.rmconsult.com/news-and-views/new-marketing-and-human-resources-professionals-join-rms-business-services-department/
8. Andrea Story's bio, on specialising in proposal development
https://www.rmconsult.com/our-people/andrea-story/
9. Your CMMC Level 2 certification announcement of December 22nd 2025
https://www.rmconsult.com/news-and-views/rm-achieves-cmmc-level-2-

certification-with-guidance-from-stratus-services/

10 Your Zweig Group 2026 placement, tenth in the 100 to 199 employee category and a fifteenth consecutive year
https://www.rmconsult.com/news-and-views/rm-named-one-of-zweig-groups-2026-best-firms-to-work-for/

11 Your What We Do page, the 24 service lines, checked for any mention of AI or automation
https://www.rmconsult.com/what-we-do/

12 Len Story's page, where we read the address this reached you at
https://www.rmconsult.com/our-people/len-story/

13 DOWL's 2020 article on big data and machine learning, which claims no deployment
https://www.dowl.com/news/our-geobusiness-future-big-data-machine-learning-and-ai/

14 DOWL's news index, page one, checked for AI and automation
https://www.dowl.com/company/news/

15 CRW Engineering's homepage, checked for AI and automation
https://www.crweng.com/

16 RESPEC's data and technology page, checked for AI and machine learning
https://www.respec.com/expertise/data-technology/

17 HDR's AI statement, on prohibiting generative AI in any final work product
https://www.hdrinc.com/artificial-intelligence-informational-statement

18 Michael Baker's April 16th 2026 announcement of its platform, its reach and its agent library
https://www.prnewswire.com/news-releases/michael-baker-international-unveils-titan-an-enterprise-ai-platform-powering-a-new-way-of-working-302745052.html

19 ASCE Policy Statement 573, adopted July 18th 2024, on where engineering responsibility sits
https://www.asce.org/advocacy/policy-statements/ps573---artificial-intelligence-and-engineering-responsibility

20 The 2026 Ames and Gough survey of fifteen A and E liability insurers, on AI adoption as a disruptor

https://www.insurancebusinessmag.com/us/news/construction/most-aande-liability-insurers-plan-rate-hikes-in-2026--ames-and-gough-566403.aspx

21 The ISO generative AI exclusion endorsements effective January 2026, and that they are optional
https://www.independentagent.com/vu_resource/verisk-to-roll-out-new-general-liability-exclusions-for-generative-ai-exposures/

22 The Louisiana DOTD evaluation of photo based field inspection, the one independent measurement in this field
https://rosap.ntl.bts.gov/view/dot/56247

23 The federal permitting drafting benchmark, published by the vendor and assessed by nineteen subject matter experts
https://www.gend.co/blog/draftnepabench-ai-federal-permitting

24 The outside view we hold ourselves to, on how often enterprise AI pilots show no measurable result
https://finance.yahoo.com/news/mit-report-95-generative-ai-105412686.html