

FIELD STUDY PREPARED FOR THE TATITLEK CORPORATION

Tatitlek already runs GovEagle, and GovEagle's own pages cover most of the build we came here to sell you.

ANCHORAGE, ALASKA SEPTEMBER 19TH, 2026

IN SHORT

We read your capture operation from the outside for a day, expecting to find you an AI build. We found something more useful, so this study says it first.

Your own requisition 2026-2731 names GovEagle as your Artificial Intelligence tool and puts your pursuit pipeline inside it. We read GovEagle's own pages and mapped the duty list in that requisition against them. The map is in the next section and it costs you nothing.

Of the eight lines we split it into, four come back covered by software you already pay for. Two are partly covered. One is unverified. One appears on none of the pages we read, and that is data calls. Separately, GovEagle's solutioning page killed the build we had planned to recommend, because it already maps task areas to your past performance.

The same requisition puts "participation in any and all data call activities as needed" on a Dahlgren seat posted at \$130,000 and up, beside a Bid and Proposal team that owns process, scheduling and compliance. Our guess, and it stays a guess until somebody measures it, is that answers on a footprint the company puts

at 1,300 employees, 46 locations and seven subsidiaries have further to travel than answers inside one office.

So what we would sell is a measurement, not a tool. One real data call measured end to end, plus the owner map that falls out of it, at eight to twelve thousand dollars. The specification is a further eight to ten and only gets written if the measurement says it should. What would get built is a written map of who owns which kind of answer, a clock, and one screen showing what is outstanding. No model anywhere in it.

The arithmetic below does not flatter us. At our central assumptions the five year build path returns about 67 percent of its cost, and 14 percent at conservative ones. Only the aggressive column pays back, at month 27. Those are recovered hours, not cash.

One thing could make all of this wrong. If GovEagle's team assignments already reaches contributors outside the proposal team, the tracker is a duplicate and it dies unbought. That is one email to your rep, it costs nothing, and if the answer is yes we will say so and stop. Send it before you reply to us.

01

Your own job posting told us you already bought an AI tool, which is not what we expected to find

We went looking for the gap in your capture operation. The first thing we found was a tool you already own that covers most of it.

Requisition 2026-2731, live on your careers board on September 19th, reads: "This role maintains an individual pipeline of assigned opportunities in our Artificial Intelligence (AI) Tool, GovEagle, and drives capture activities on those pursuits, and contributes to RFP response efforts as needed, working in close coordination with the Bid and Proposal team."

That is an AI capture tool already in production, holding assigned pursuits from identification through award decision. So we read what GovEagle actually does.

GovEagle

named in your own requisition as your Artificial Intelligence tool, holding assigned pursuits from identification through award decision.

The map below is that reading, duty by duty. GovEagle's solutioning page is what killed the build we had already scored as the strongest thing we could make for you, because it says it "maps task areas to your past performance and shows where you are strong, where you have gaps", and its own writing says "The AI drafting layer pulls from past proposals, past-performance write-ups, and boilerplate repositories to generate first drafts". The last two rows of the map are the ones that point at what is left.

The coverage map. The duty list on requisition 2026-2731 as we read it on September 19th, split into eight lines, against what GovEagle's own pages say. This is the free part, and it is the reason the rest of this study is short.

The duty, from requisition 2026-2731	What we found, and on which GovEagle page	Verdict
Maintain a pipeline of assigned opportunities, logging, staging and updating each pursuit from identification through award decision	Opportunity tracking, pursuit qualification and pipeline analytics, on the platform page	Covered
Conduct capture activities leading up to and through RFP submission	RFP shredding, compliance matrix and outline generation, on the platform page	Covered
Develop capture plans and gate review materials for assigned pursuits	Capture deck, solution brainstorm and opportunity fit score, on the solutioning page	Covered
Coordinate with operations, contracts, pricing and program staff to align capture strategy with execution capacity	Bid and no-bid analysis, task area matching and early go or no go guidance, on the solutioning page. Nothing we read says it reaches your operations, contracts or pricing staff	Partly, the go or no go half is in the tool and the coordination half is not

The left column compresses the posting's own duty list, and the full wording is in the posting. Two of the posting's sentences carry two duties each, so rows two and eight are the two halves of one sentence and rows three and six are the two halves of another. The four pages read on September 19th were [goveagle.com/platform](#), [/solutions/solutioning](#), [/solutions/proposals](#) and the automation-software post. That is a bounded search and not proof of absence, which is why the last two rows say ask rather than no. Every Covered row is a line an AI vendor could still quote you for. Mark each row confirmed, corrected or unknown with whoever administers GovEagle, and the corrections are worth more to us than the agreements.

The duty, from requisition 2026-2731	What we found, and on which GovEagle page	Verdict
Partner with the Bid and Proposal team through the response effort	Color Team reviews and team assignments, on the platform page. The review and assignment mechanism sits in the tool, the partnering is still people	Covered
Develop price to win analysis	A pWin calculator on the platform page. Whether it produces the price to win number itself is not stated	Partly, and we would leave this one human either way
Support the development and scrubbing of key personnel resumes, so they align with solicitation requirements	Resumes are named on none of the four pages we read. Search summaries assert the feature, which is not a page we fetched	Unverified, ask your rep
Participate in any and all data call activities as needed	Data calls appear on none of the four pages we read	Not found, and it is the reason for the rest of this study

The left column compresses the posting's own duty list, and the full wording is in the posting. Two of the posting's sentences carry two duties each, so rows two and eight are the two halves of one sentence and rows three and six are the two halves of another. The four pages read on September 19th were goveagle.com/platform, /solutions/solutioning, /solutions/proposals and the automation-software post. That is a bounded search and not proof of absence, which is why the last two rows say ask rather than no. Every Covered row is a line an AI vendor could still quote you for. Mark each row confirmed, corrected or unknown with whoever administers GovEagle, and the corrections are worth more to us than the agreements.

The mistake here is not moving slowly, it is buying the same capability twice

We can't tell you what your capture operation costs, because none of it is published. We can tell you what we nearly sold you.

You already carry a line for a capture platform. The ordinary next line beside it is a past-performance tool, and GovEagle's own pages say it already does that work. What such a tool costs is not published anywhere we could reach, so we can't price the mistake, only name it.

In our experience the person buying the second tool is rarely the person who read the first tool's feature list a year ago.

Four of eight

lines of your own duty list that GovEagle's pages cover outright, with two more partly covered. The map is free and it is in the section above.

There is a second cost and we can't price it either. The Dahlgren seat is live now, and a shared method costs less to land beside a new hire than on top of one who has built their own. That is a belief rather than a finding.

03

What changes is where a proposal manager looks to find out what is still outstanding

A data call walked step by step.

A solicitation lands and the compliance matrix comes out of GovEagle. Its items need answers that may sit with any of the subsidiaries. Here is the ordinary shape of that

work, and what it looks like against a completeness board.

One data call, from the matrix to the last answer in

THE ORDINARY SHAPE	WITH A BOARD
1 The compliance matrix comes out of GovEagle.	Same matrix, same tool, imported as a file.
2 Each item has to be matched to whoever owns the answer, by asking around until somebody claims it.	A written owner map resolves most items by rule. Every miss is listed by name.
3 Who has answered is tracked wherever the coordinator keeps it.	The board holds the list, and the proposal manager approved who was on it before anything sent.
4 Chasing is remembered by hand, alongside everything else in the week.	Chases go on a cadence a person set, from your own mail domain, with delivery receipts.
5 The outstanding count is assembled when somebody needs it.	Open the board and read the count.
6 A request that never arrived looks the same as one nobody has answered yet.	A silent chase shows as a delivery failure rather than as a non-response.
7 The named owner decides whether the answer they gave is correct and complete.	The named owner decides whether the answer they gave is correct and complete.

The left column is how this work generally runs. It is NOT a claim about how yours runs, because we can't see that from outside, and replacing it with your own numbers is exactly what the measurement is for. Row one names GovEagle and is about you. Row two is struck on neither side because a written map resolves most of that step and not all of it. The last row is identical on purpose: whether a returned answer is correct is a judgement with somebody's name on it, and nothing we would build touches it. Even with every struck row gone, our conservative column still assumes only a quarter of coordination time comes back.

What the columns can't carry is the part that decides whether any of this is worth doing. Nothing we read tells us your data calls lose time. The whole evidence base is one clause in one job posting. So the first thing we would do is measure one.

04


The thing we would build has no AI in it, and saying so is the point

You already have the AI. What is missing is a list, a clock and a board.

A written map of who owns which kind of answer across the subsidiaries. A list of what a data call needs, from the matrix GovEagle already makes. Due dates. Reminders from your own mail domain. An escalation step. One screen showing what is outstanding. A log of who was asked, when, and what came back.

The matrix comes in as a file, and if GovEagle gives you only a screen, that first step becomes a paste and nothing else changes. Rules match each item to a named owner from the map. The proposal manager approves the recipient list before anything sends, because nothing should mail people across seven subsidiaries without a person saying yes. Chases go on a cadence a human set. Contributors mark an item submitted and attach a pointer to their own file.

It tracks status. It never writes into a submitted volume, never decides whether an answer is true, and never generates prose.



Three built pieces, one list, and no model anywhere in the path. GovEagle stays the source of what a data call needs. The tracker only holds status against it, and every piece is one small system GeoNorth can host and own.

A working demonstration, honest about being one. [Open it full screen.](#)

Nothing in the design we would ship needs a model to make a decision. Rules can be tested exhaustively and give the same answer every time, which is what a document heading toward a contracting officer needs. Chain four judgement steps instead, each right nine times in ten, and the whole run is right about two thirds of the time, which across a real matrix means several silently wrong rows on every pursuit with nobody assigned to catch them.

One place could ever earn a single model call, mapping a free-text requirement to an owner when no rule matches. We are not proposing it, and the specification counts the misses so that count has to justify it later. The price to win number stays human too, because the cost of a wrong one is an award, and that duty is named in your own requisition.

A contractor has alleged in the Court of Federal Claims that the Army used AI that hallucinated during a source selection, and the Army conceded one weakness it had identified was not supported by the record. GAO denied the related protest in May 2026. That is the government's tooling rather than a contractor's, so it is no precedent for anything you would run. It is why nothing we build authors a fact.

The logins, the mail and the database are bought. The owner map and the routing spine are built, because no product knows that a given question belongs to a named person at a named subsidiary.

We have no way to see whether that map exists as one artefact today, and in the operations we have worked in it usually does not. Assembling it from one live pursuit is most of what the fee buys, and it goes stale faster than it goes wrong, which is why the roadmap checks it again at 90 days.

If the measurement shows a low item count and few chases, the right answer is a bought list, the owner map and a written cadence, and no software.

GeoNorth is your own information systems subsidiary, which makes it the obvious place for this to live, and whether it has capacity is your call. There is no licence and no hosting invoice from us.

05

The build does not pay back under our conservative or our most likely assumptions, so we are not asking you to fund the build

Every driver below is an assumption, because nothing about your data calls is published and we refuse to invent it.

Three scenarios on the full five year path, our fee plus a tracker built and run inside GeoNorth. We ran the arithmetic in code, so you can redo it. The conservative column recovers 14 percent of its five year cost, the most likely 67 percent, and only the aggressive column pays back, at month 27. Those are recovered hours, not cash.

The measurement is eight to twelve thousand dollars against a five year commitment this table puts between ninety and a hundred and twenty seven thousand. If it kills the build, the money bought the kill. The driver that moves the answer most is your annual data call count, and that is a number you can pull off your own pipeline for nothing.

The full five year path, three scenarios. The provenance mark on each row says where the number came from.

	Conservative	Most likely	Aggressive
☒ Posted floor on the seat that carries this work	\$130,000	\$130,000	\$130,000
☐ Data calls a year	10	20	32
☐ Coordination hours a data call	8	12	16
☐ Loaded hourly rate used	\$62	\$78	\$88
☐ Share of coordination time removed	25%	35%	45%
☒ Hours a year recovered at full run rate, all three labour lines	43	204	538
☒ Benefit a year at full run rate	\$4,151	\$17,292	\$47,069
☒ Five year cost of ownership	\$126,816	\$111,696	\$90,643
☒ Share of that cost recovered in five years	14%	67%	228%
☒ Pays back inside five years	No	No	Month 27

The assumptions, in full. Data call volume of 10, 20 and 32 a year is assumed and is the driver that moves the answer most; nothing published states your pursuit or data call count. Hourly rates of \$62, \$78 and \$88 are assumed for capture and proposal time, anchored on the \$130,000 posted floor in requisition 2026-2731, which is a floor on one seat and not a fully loaded cost, and the conservative column adds no employer burden on top of it. Contributor time is priced separately and lower, at \$55, \$62 and \$68, because the people answering a data call are not capture staff. Coordination hours of 8, 12 and 16 a data call are assumed, and that is the one labour driver the pilot actually measures, as are 3, 5 and 6 hours finding who owns an answer and 8, 14 and 20 aggregate contributor hours. The recovered hours are about one working week a year conservative, about five most likely and about thirteen aggressive. Cut percentages are assumed throughout, at 25/35/45 percent on coordination, 35/50/60 on finding the owner and 15/25/30 on contributor time. One avoided-cost line sits in the benefit row beside those three: one rebuild every five years of a tracker specified without a baseline, at an assumed \$20,000, \$25,000 and \$30,000 a rebuild, of which 40, 50 and 60 percent is avoided by holding a measured baseline first. That line is the only place this table counts the measurement's own value. The five year cost carries an internal build of \$48,000, \$44,000 and \$36,000 assuming GeoNorth builds it, our fee at \$22,000, \$18,000 and \$16,000, split into the measurement and owner map at \$8,000 to \$12,000 and the specification at \$8,000 to \$10,000, training of 140, 110 and 90 hours priced at each column's own rate, run cost of \$6,000, \$5,000 and \$4,200 a year over 4.5 years, and contingency of 20, 20 and 15 percent. There is no token or inference line because there is no model in the shipped path. Year one runs at 25, 30 and 40 percent of run rate for a mid-year go live and about ninety days of stabilisation, and years three to five are flat because a completeness board has no mechanism that makes it better on its own. Deltek's reported average of 84 hours to develop a proposal comes from a survey of 917 contractors and is not a measurement of you, so it sets no driver here. All three labour lines are capacity, not cash, and they only become money if the recovered hours visibly go back into capture plans, price to win work and gate review material.

☒ verified against a source ☒ modelled from stated drivers

☐ assumed, and named as such

14%

What the conservative case recovers of the full five year cost, on about 43 hours a year, which is a little over one work week. The most likely case recovers 67 percent and still does not pay back. These are recovered hours and not cash, and they only become money if the hours visibly go back into capture plans and price to win work. We are showing you the column that argues against the build because it is the one that decides the question.

The outside view, from your own market. GovConWire, reporting on Deltek's 2026 study of 917 government contractors, put 90 percent of them using AI in some capacity, and Deltek's own release on the same study reported 45 percent unclear on the return. Nearly everyone here is adopting and a large share can't tell what they got back. We have no published failure data for the class this build belongs to, so treat this as our judgement: in coordination tooling, what kills it is a stale owner map, not a bad model. The nearest published result to what we describe is a cloud vendor's write-up of Koniag Government Services, where a narrow automation with no model in it cut a one hour task to two minutes. That is a vendor publishing a partner's result rather than an audit, so weigh it as such.

The number sits with you today, and moves to the Dahlgren capture manager once that seat is filled. The proposal manager holds the measurements, the Bid and Proposal team holds the completeness data. Be clear what clearing means: below the aggressive column on the hours alone, the honest call is to stop, because the most likely column does not pay back. Two of the four drivers, your volume and the hours a call takes, are known at gate one, so we re-run this table in code then and publish which column you landed in, on your own loaded rate if you give us one and on ours if not. The fourth, how much time a board really removes, is only visible at gate three after the build money is spent.

06

Start with the free part, because it can end this before you spend anything

Sequenced against your calendar, not ours.

The coverage map is above and it is yours. Three questions come out of it and only your GovEagle rep can answer them: whether team assignments reaches contributors outside the proposal team, whether it handles key personnel resumes, and whether the compliance matrix exports as a file. The first can end this before it starts, because if team assignments already reaches those contributors the tracker is a duplicate and it dies there, unbought.

The measurement runs on the first suitable pursuit after September 30th. Measuring inside the federal year end rush would get in the way, and would produce a baseline you could fairly dismiss as atypical.

Phase	What runs	How we would know it worked
Now	Take the coverage map in the section above and mark each line confirmed, corrected or unknown with whoever administers GovEagle. Put the three open questions to your account rep. Separately, count the data calls you ran in the last twelve months off your own pipeline, which costs you an hour and no fee.	Every line of the map marked confirmed, corrected or unknown, including the two we scored partly, each carrying the page it came from. All three rep questions answered in writing. One number for annual data call volume, from your records.
	Measure one real data call end to end, on one live pursuit, after September 30th. Change nothing about how it runs.	Six numbers, measured on one real call: elapsed hours from request to last answer, distinct people asked, chase rounds, items where the first person asked was wrong, items outstanding at your internal cutoff, and proposal manager hours spent routing and chasing. With your own volume count from gate zero, that is the seven the decision needs.
	Extract the owner map from that same pursuit and write down the method, so your team can repeat it without us.	Share of the matrix items whose owner is resolved by a written rule rather than by asking. Our bar is 80 percent, and below it we report that rules are the wrong spine.
Next	The tracker specification with its acceptance bar, written in the six measured numbers, for you to hand to GeoNorth.	GeoNorth can price and schedule the build from it without a question back to us.
	Run the built tracker in the shadow of your current method on one real data call, before it becomes anyone's method.	Coordination hours, the proposal manager time spent routing and chasing, measured against the gate one baseline. That is the funding test, because it is the number the ROI runs on. A fall below the conservative column's 25 percent stops the spend.

Phase	What runs	How we would know it worked
		Between 25 and 45 percent the build does not pay back on hours, and we report that in writing rather than treating the gate as passed. Items outstanding at cutoff and re-routes are reported beside it and are not the funding test.
	Check the owner map for decay at 90 days against a live pursuit, because our design assumption is that this class of tool dies of a stale map rather than of a wrong one.	Share of owner map entries still correct after 90 days. Our bar is 90 percent, and below it the map needs a named owner and a refresh cadence before the tracker is trusted.
Later	A register of which AI tools are in use, who operates each one and which contract it touches, with GovEagle as row one. It is a form and a table and an owner, and it would be dishonest to call it an AI build. Washington Technology reports that GSA posted an AI acquisition rule to the Federal Register on June 17th, 2026 whose 120 day disclosure and 72 hour incident clauses are proposed and not in force, and whether your capture use falls inside them is unsettled, since they address models used in the performance of a contract and capture sits before award. Worth having on its own merits.	Time from a question arriving to a complete answer about which tools are in use and where. Unscoped until somebody asks for it, so there is no bar on it yet.
	A written table of the five certifications and vehicles you publish, ISO 9001:2015, ANC 8(a) sole source, SEAPORT NXG, GSA STARS III and OASIS+, with the ordering path and reporting for each, so a capture manager stops asking. It is a table before it is anything cleverer, and it only earns a search layer if the table turns out	Time for a capture manager to get the correct ordering path for a named vehicle. Unscoped until somebody asks for it, so there is no bar on it yet.

Phase	What runs	How we would know it worked
	to be too long to read, which with five entries it will not.	
	Refused, and listed so nobody circles back to it. A bid and no-bid capacity check, because GovEagle already ships bid and no-bid analysis, an opportunity fit score and pWin, and the rest lives inside your own staffing and pricing systems where GeoNorth is better placed than we are.	None. We will not do this work.
	Refused. A key personnel and past-performance engine, which was our own highest scoring candidate until we read GovEagle's pages. Its past-performance half duplicates what you already pay for. Its resume half is unverified in both directions and stays a question for your rep, not a build.	None. We will not do this work.

Four gates and each one is a real way out. Gate zero costs nothing and happens before you buy anything: the coverage map, the three rep answers and your own volume count. A yes on the team assignments question kills the tracker unbought. If your rep will not answer in writing, treat the row as unverified rather than as a no, and the measurement still runs. What we will not do is build against a feature we assumed away. Gate one is the honest one. Before the measurement runs, your proposal manager and BD lead write down what a process that holds looks like, on those six measured numbers plus your volume count. Then we measure and compare against your bar, not ours. If it clears, the specification does not get written and does not get billed, which is eight to ten thousand dollars you keep, and we say in writing that you should buy nothing further from us. Gate one also re-runs the cost table in code on your measured volume and hours, and publishes which column you landed in. Below the aggressive column the build is not justified by hours alone. Gate two is build versus buy and you should expect us to lose it. Gate three is the pilot: the tracker shadows your current method, and the funding test is coordination hours rather than the outstanding count, which is the easier number to move. Below 25 percent the spend stops. Between 25 and 45 percent we write down that the build does not pay back on hours. Note what gate three does not settle: it measures the coordination driver only, so the owner-finding and contributor lines stay assumed, and together they are the larger share of the modelled benefit at every column.

One live pursuit we can measure after September 30th, with your proposal manager's agreement. That is the single dependency that moves the timeline. Forty five minutes with whoever administers GovEagle, to mark the coverage map. The three questions put to your rep. One real compliance matrix export, or a screenshot. A named proposal manager as decision owner, and a contact at GeoNorth if GeoNorth is to build it.

07

Four things could make this study wrong, and the first one would end it

In the order they would actually bite.

GovEagle may already do this. Team assignments is named on its platform page, and our reading that it sits inside a pursuit rather than reaching contributors across your subsidiaries is an inference we drew, not something we read. It is the thing most likely to make this wrong. The answer costs one email, the map that surfaces it is free, and if the answer is yes we will say so and stop. We would rather be wrong for nothing than right for a fee.

Your data calls may be fine. Nothing we read says they are slow or expensive, and we will not imply otherwise to make a sale. One clause in one job posting is the entire basis for thinking there is a problem here. That is why the first paid thing is a measurement, and why gate one is written so a good result ends the engagement.

We have not seen GovEagle export a matrix as a file. No page we read says it does, and the tracker's first step assumes it. If it gives you only a screen, the intake becomes a paste and nothing else changes, so this costs a step rather than the project.

The numbers are assumptions and the rail says so on every row. One figure is verified and it is a posted salary floor. One call is one call, so we would report the range across its items rather than an average, and a second measured call is the cheapest way to learn whether the first was typical.

The next step is one email you send to somebody else

Ask your GovEagle rep three things. Whether team assignments reaches contributors outside the proposal team. Whether it handles key personnel resumes. Whether the compliance matrix exports as a file. Those answers decide whether there is anything here worth paying for, and we can't get them. Tell us what comes back either way.

What we checked

Every number in this study traces to a page we read. If we have any of it wrong, that is worth knowing too.

- 1** Requisition 2026-2731, a Strategic Proposal Support and Capture Manager seat at Dahlgren, Virginia posted at \$130,000 and up, names GovEagle as the company's Artificial Intelligence tool, puts the pursuit pipeline in it, and lists the data call and resume duties
https://careers-tatitlek.icims.com/jobs/2731/strategic-proposal-support-%26-capture-manager/job?in_iframe=1
- 2** GovEagle's platform page names RFP shredding, the compliance matrix, outline generation, Color Team reviews, the capability matrix, pWin, pipeline analytics and team assignments
<https://www.goveagle.com/platform>
- 3** GovEagle maps task areas to past performance and matches PWS requirements to your strongest projects
<https://www.goveagle.com/solutions/solutioning>
- 4** GovEagle's AI drafting layer pulls from past proposals, past-performance write-ups and boilerplate repositories

<https://www.goveagle.com/blog/govcon-proposal-automation-software>

- 5 GovEagle extracts requirements, organises them by section and maps them to the response outline
<https://www.goveagle.com/solutions/proposals>
- 6 1,300 employees across offices throughout Alaska and 46 locations worldwide, the seven subsidiaries, and the published certifications and vehicles: ISO 9001:2015, ANC 8(a) sole source, SEAPORT NXG, GSA STARS III and OASIS+
<https://tatitlek.com/our-company/>
- 7 More than 300 shareholders
<https://tatitlek.com/>
- 8 Mark Wamsher, Vice President of Business Development
<https://tatitlek.com/leadership/>
- 9 Deltek's 2026 GovCon Clarity study surveyed 917 contractors, and 45 percent reported being unclear on the return on AI
<https://www.deltek.com/company/news/latest-deltek-clarity-industry-studies-highlight-ai-challenges/>
- 10 Deltek reported 90 percent of contractors using AI in some capacity, up from 45 percent, and an average of 84 hours to develop a proposal
<https://www.govconwire.com/articles/clarity-report-2026-deltek-govcon-kevin-plexico>
- 11 A contractor alleges the Army used AI that hallucinated during a source selection, the Army conceded one identified weakness was not supported by the record, and GAO denied the related protest in May 2026
<https://www.nextgov.com/acquisition/2026/08/contractor-alleges-army-inappropriately-used-ai-make-450m-contract-award/415225/>
- 12 GSA posted an AI-specific acquisition rule to the Federal Register on June 17th, 2026, with proposed 120 day disclosure and 72 hour incident notification
<https://www.washingtontechnology.com/contracts/2026/06/gsa-unveils-new-ai-specific-acquisition-rule/414351/>
- 13 A cloud vendor's case study credits Koniag Government Services with cutting a one hour vendor responsibility task to two minutes
<https://aws.amazon.com/solutions/case-studies/us-army-koniag-gov-services/>
- 14 GeoNorth Information Systems is the group's information systems subsidiary
<https://tatitlek.com/geonorth-information-systems-llc/>

